

Lecture 9: Reductions

Mark Saroufim

Follow along

Chapter 10 of PMPP book

Locally or remotely <https://lightning.ai/>

git clone <https://github.com/cuda-mode/lectures>

cd lecture9

nvcc -o sum *.cu

ncu sum

What's a reduction

Operations that reduce the output size

Most typical take a vector and produce a scalar

min, max, argmax, argmin norm, sum, prod, mean, unique

Demo: `torch_reductions.py`

Reductions are everywhere

- Mean/Max pooling
- Classification: Argmax
- Loss calculations
- Softmax normalization

Reductions in PyTorch

<https://github.com/pytorch/pytorch/blob/main/aten/src/ATen/native/cuda/ReduceOps.cpp>

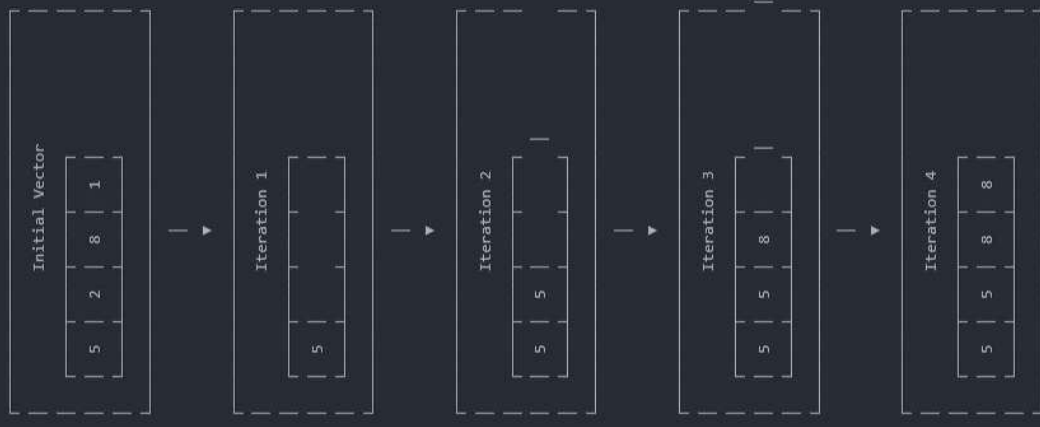
```
>>> a = torch.randn(1, 3)
>>> a
tensor([[ 0.6763,  0.7445, -2.2369]])
>>> torch.max(a)
tensor(0.7445)
```

Serial reduction example

Max operation

Go through elements 1 by 1

Compare new number to old max if greater then update



More general formulation

```
def reduce(data, identity, op):
    result = identity
    for element in data:
        result = op(result, element)
    return result

# Example usage:

# Summation
data = [1, 2, 3, 4, 5]
print(reduce(data, 0, lambda a, b: a + b)) # Output: 15

# Product
print(reduce(data, 1, lambda a, b: a * b)) # Output: 120

# Maximum
print(reduce(data, float('-inf'), max)) # Output: 5


# Minimum
print(reduce(data, float('inf'), min)) # Output: 1
```

<https://gist.github.com/msaroufim/a062aa0b08a4cc57e02db634a67c6b20>

Transformation vs reduction

What should the thread strategy be?


Output size < Input size that's why we call them reductions



Transformation:

e.g. $c[i] = a[i] + 10;$

Thread strategy: one thread per output point



Reduction:

e.g. $*c = \sum a[i]$

Thread strategy: ??

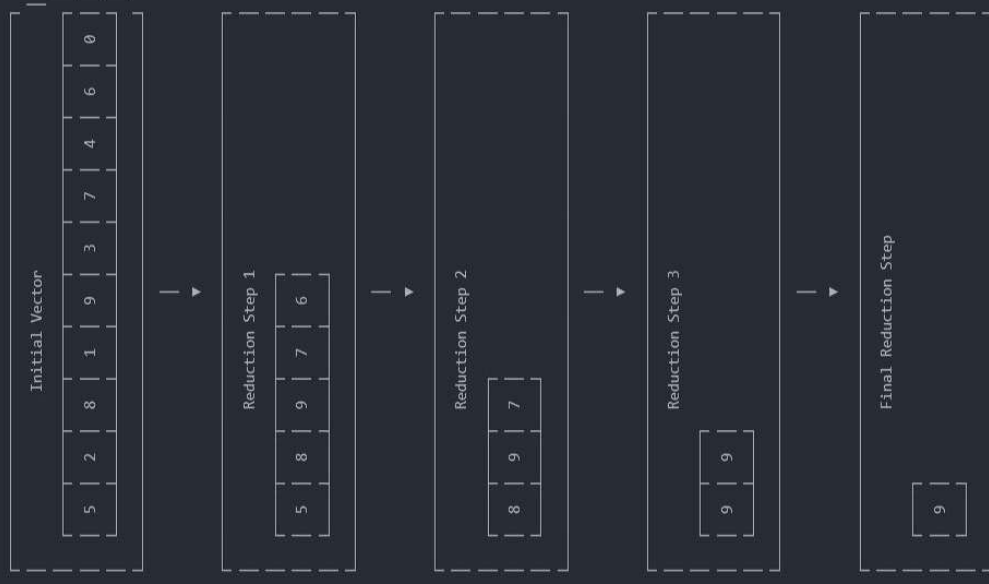
<https://www.youtube.com/watch?v=D4I1YMsGNIU&t=1763s>

Parallel Reduction visualization

At each step take a pair of elements and compute their max and store the new max in new vector

Continue until there is 1 element in the vector

$O(\log n)$ steps



Reduction Trees:




FIGURE 10.5

A parallel sum reduction tree.

Non determinism and accuracy

`torch.use_deterministic_algorithms(True)`

Demo

- `nondeterminism.py`
- `accuracy.py`

Reduction Kernel




FIGURE 10.7

The assignment of threads (“owners”) to the input array locations and progress of execution over time for the SimpleSumReductionKernel in Fig. 10.6. The time progresses from top to bottom, and each level corresponds to one iteration of the for-loop.

- A lot of threads will be inactive :(
- A lot of warps (groups of 32 threads) will be inactive :(
- Let's check ncu -set full

simple_reduce.cu

Remember the performance checklist

Lecture 8!

- Control divergence
- Memory divergence
- Minimize global memory access
- Thread coarsening

Minimize Control Divergence

Ensure threads and their owned positions remain close together as time progresses

Quiz: Which other problem does this fix?




FIGURE 10.8

A better assignment of threads to input array locations for reduced control divergence.

control_divergence_reduce

Minimize Global Memory Access




FIGURE 10.10

Using shared memory to reduce accesses to the global memory.

shared_reduce.cu

Hierarchical reduction

Let's try running input size 4096




FIGURE 10.12

Segmented multiblock reduction using atomic operations.

`segment_reduce.cu`

Thread Coarsening (Andreas' favorite optimization)




FIGURE 10.14

Thread coarsening in reduction.

`reduce_coarsening.cu`

Next steps

Lecture 1-8 gave you everything you need to start writing, profiling and shipping kernels in PyTorch so start picking a project - Look for collaborators in #general to stay motivated

Next Lecturer is Oscar who will talk about shipping production CUDA libraries

Looking for lecturers interested in covering prefix sum (scan) and NCCCL

Bonus slides: Reductions in the real world

Example of reductions

User facing ops

How reductions are implemented in PyTorch

- <https://github.com/pytorch/pytorch/blob/4b494d075093096d822b9d614e2719a0e821c6af/aten/src/ATen/native/cuda/ReduceMaxValuesKernel.cu#L53>
- <https://github.com/pytorch/pytorch/blob/main/aten/src/ATen/native/cuda/ReduceMaxValuesKernel.cu>
- <https://github.com/pytorch/pytorch/blob/main/aten/src/ATen/native/metal/ops/MetalReduce.mm>
- CPP style of CUDA (Might need its own lecture)

Key ideas

- Implementation is accumulator and reduction op agnostic
- TensorIterator to iterate over tensor elements
- ReduceConfig: Has kernel launch parameters like block size and number of threads, grid etc.. and its set in setReduceConfig
- Reduce_kernel is where it gets launched
- Reduction strategies: thread level, block level x,y, or global reduce
- Vectorization: Over input and/or output

torch.compile!

To the notebook - [reduce_compile.py](#)

Look out for

- ReductionHint
- tl.sum
- triton_heuristics

Triton

<https://github.com/openai/triton/blob/main/lib/Conversion/TritonGPUToLLVM/ReduceOpToLLVM.cpp>

```
// First reduce all the values along axis within each thread.  
  
reduceWithinThreads(helper, srcValues, accs, indices, rewriter);  
  
// Then reduce across threads within a warp.  
  
reduceWithinWarps(helper, accs, rewriter);
```