



Lecture 10: Building production ready CUDA libraries

A computer architecture aware programming perspective

Oscar Amoros Huguet

About me

BSC

- 3D convolutions in OpenCL, CUDA and OmpSs

UB

- Teaching assistant 9 years teaching OpenCL/CUDA and other CS topics
- 2 papers (OpenCL)
- First OpenSource library: SimpleOpenCL

StreamHPC

- Porting Boost Compute (OpenCL) to HSA and HAS-IL

AutomaticTV
2016 – Today

- AutomaticTV: soft real time, C++ and CUDA
- New OpenSource libraries cvGPUSpeedup and the Fused Kernel library -> paper on the way!

Lecture overview

Summary of already
seen concepts



Main story: why would
a “CUDA ninja” want to
create libraries?

Use cases:

- GPU communication utility
- GPU kernel libraries

Summary of already seen concepts

cudaMemcpyAsync -> D2H, H2D, D2D

CUDA Streams

GPU DRAM latencies and latency hiding

Memory bound and Compute bound Kernels

Kernel Fusion: Vertical



Lecture overview

Summary of already
seen concepts

Main story: libraries




Use cases:

- GPU communication utility
- GPU kernel libraries

Main story: 2008 approach

CPU
application




GPU
application

Main story: most of my C++/CUDA exp.



<https://automatic.tv/>



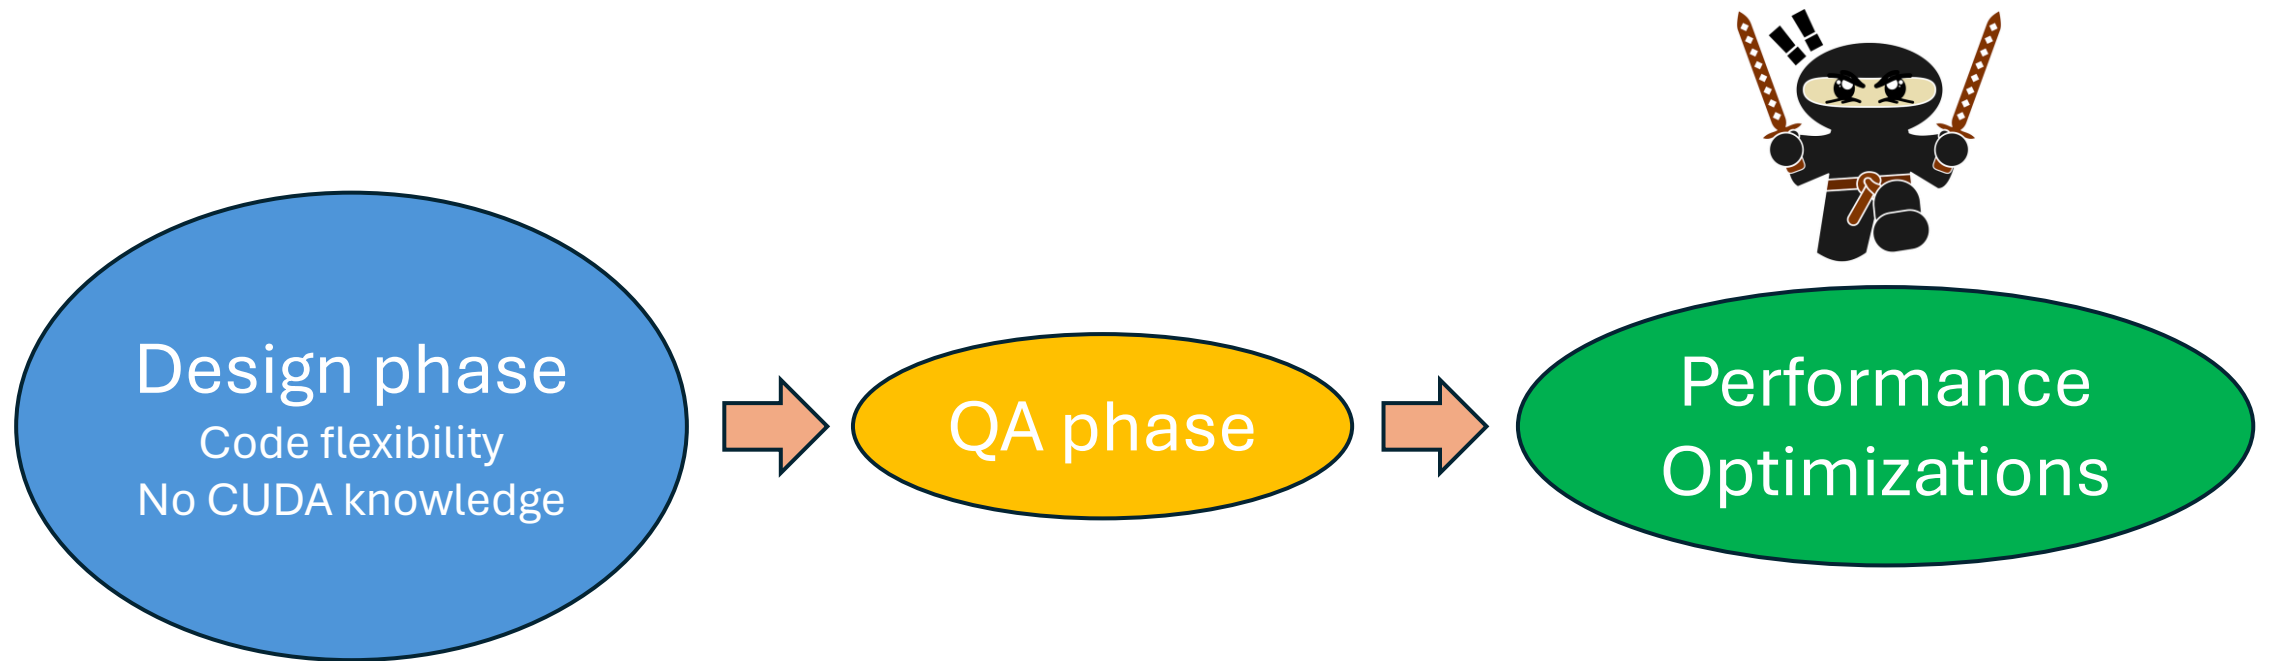
4K

4K

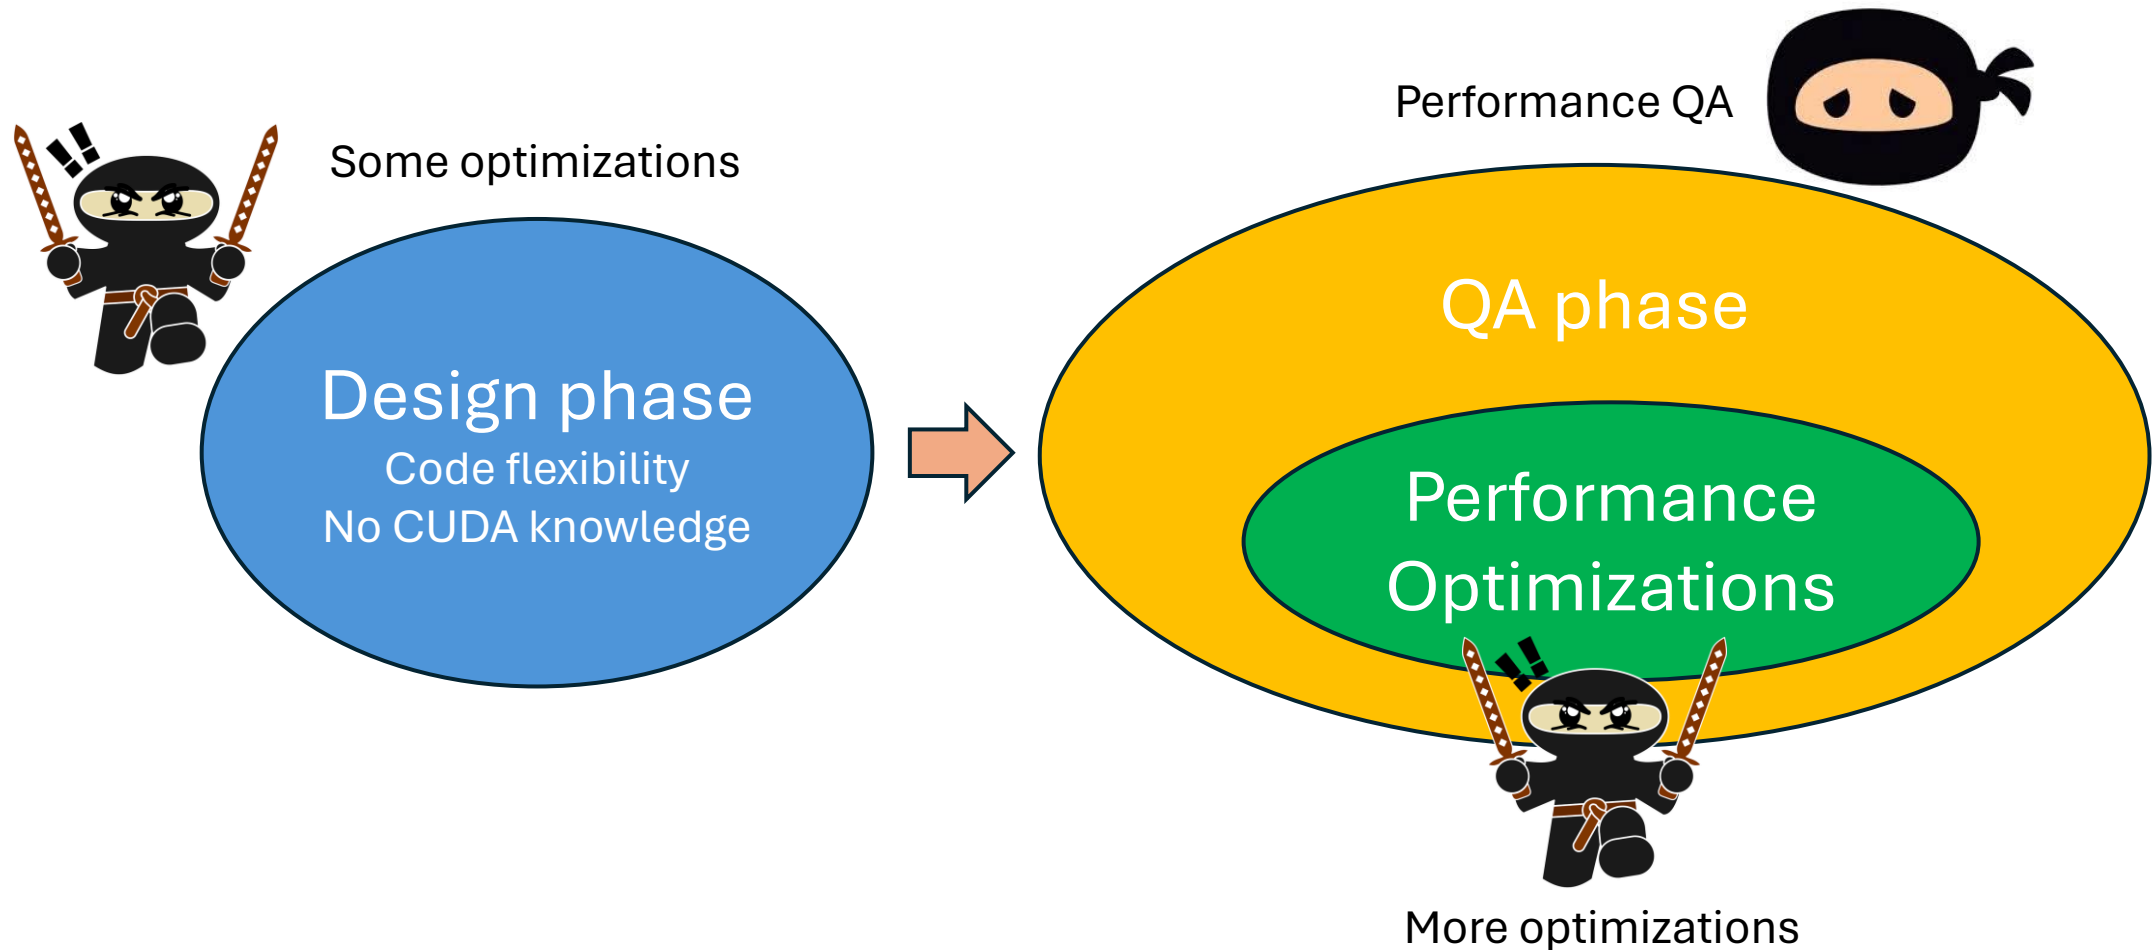


FullHD


Main story: ideal project organization



Main story: reality



Main story: problem



Main story: tough balance


Abstraction



Performance

Main story: user requirements

Abstraction



Performance

Main story: what users care about




Learning curve



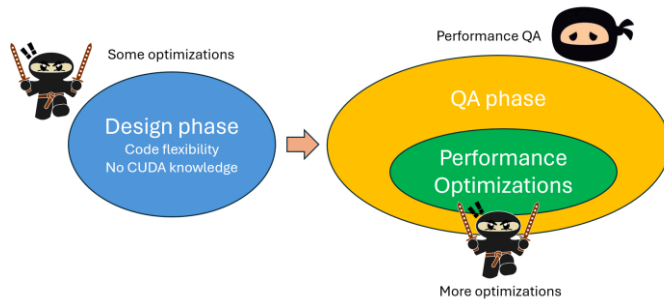
API

Main story: libraries challenges




Main story: summary


Motivation



Problem to solve



Challenges



Lecture overview

Summary of already
seen concepts

Main story: libraries

Use cases:

- GPU communication utility
- GPU kernel libraries



GPU communication utility




GPU communication utility




GPU

Memory Space Types



CPU
Pinned



CPU

GPU communication utility

Pre-allocate all your memory

GPU communication utility

Source Memory Space

!=

Destination Memory Space



```
Manager.manage(sourcePtr, destinationPtr);
```

Source Memory Space

==

Destination Memory Space



GPU communication utility

| Source MS Type | Destination MS Type | Number of copies |
|----------------|---------------------|------------------------|
| CPU | CPU | 0 |
| CPU | CPU Pinned | 1 |
| CPU | GPU | 2 CPU -> Pinned -> GPU |
| CPU Pinned | CPU | 1 |
| CPU Pinned | CPU Pinned | 0 |
| CPU Pinned | GPU | 1 |
| GPU | CPU | 2 GPU -> Pinned -> CPU |
| GPU | CPU Pinned | 1 |
| GPU | GPU | 0 |
| GPU | Peer GPU | 1 |
| GPU | Fake Peer GPU | 2 GPU -> Pinned -> GPU |

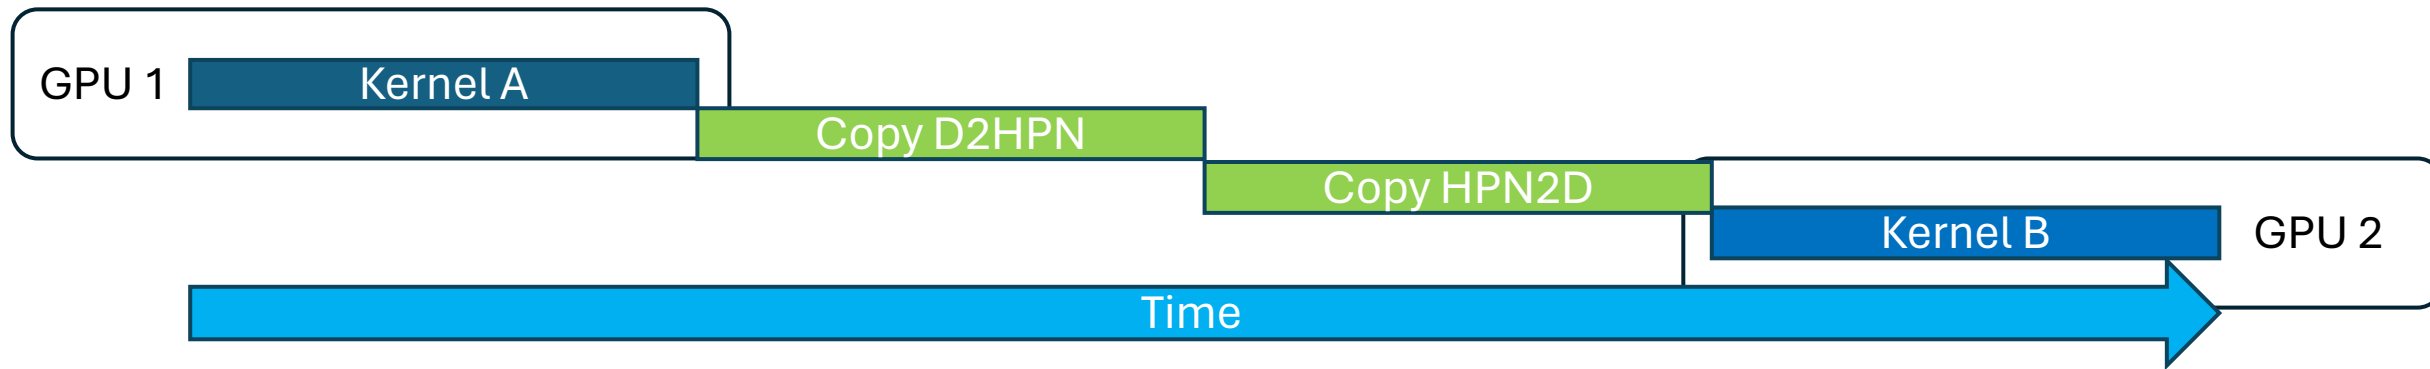
Multi GPU

cudaMemcpyPeerAsync

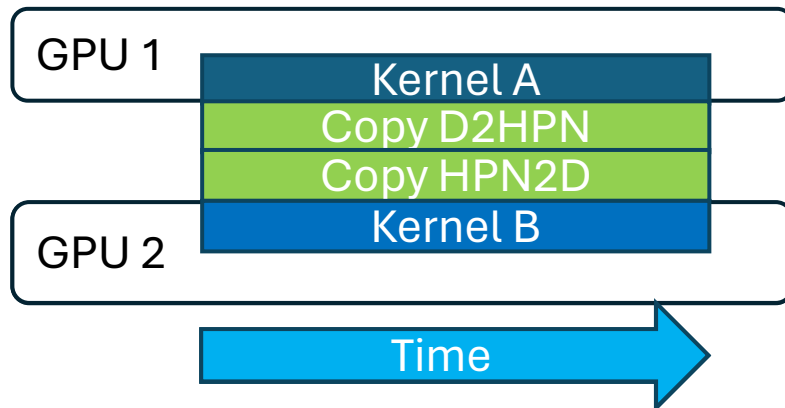
cudaMemcpyPeerAsync

GPU communication utility

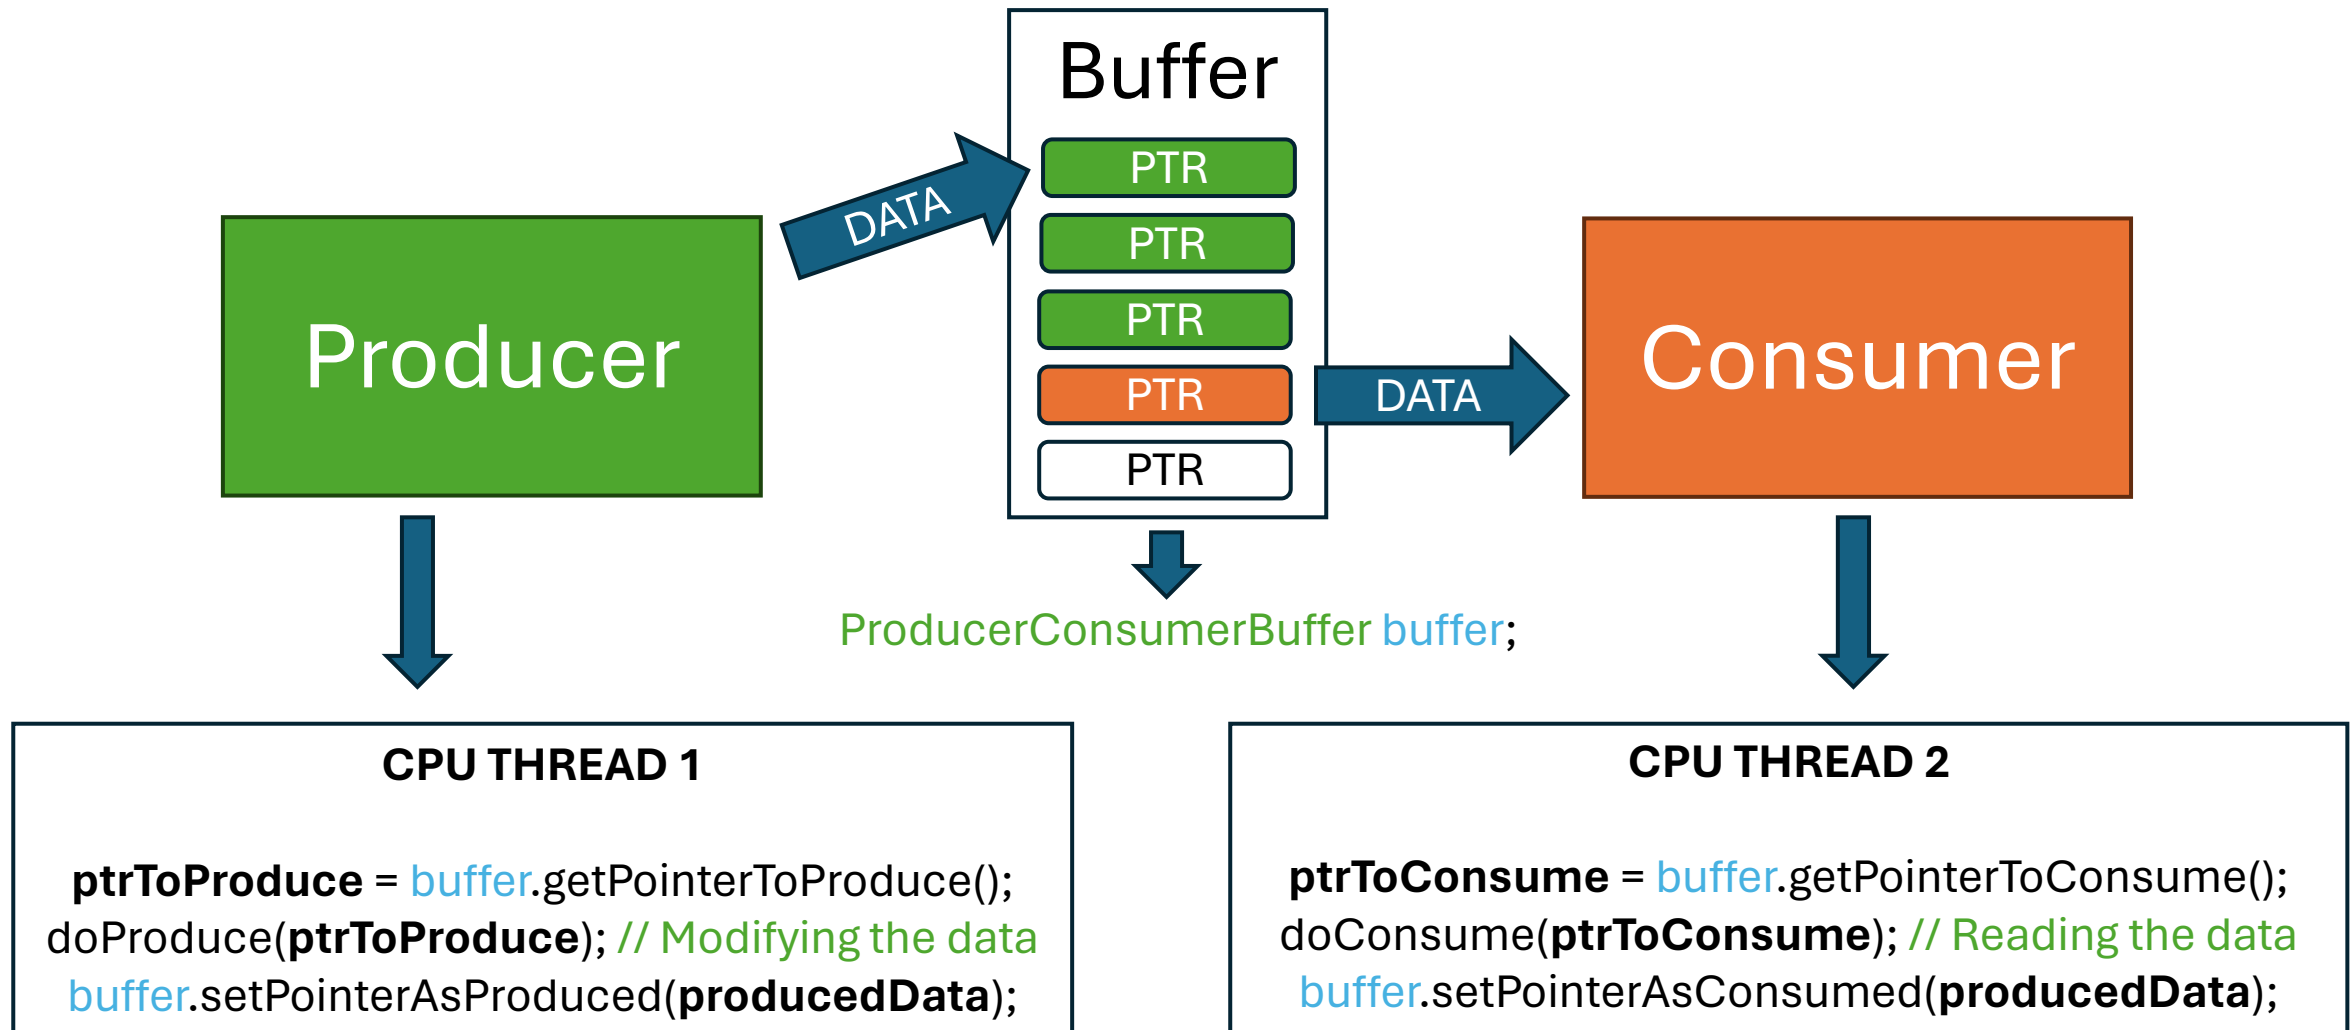
Problem: serial execution (GPU to Fake Peer GPU)



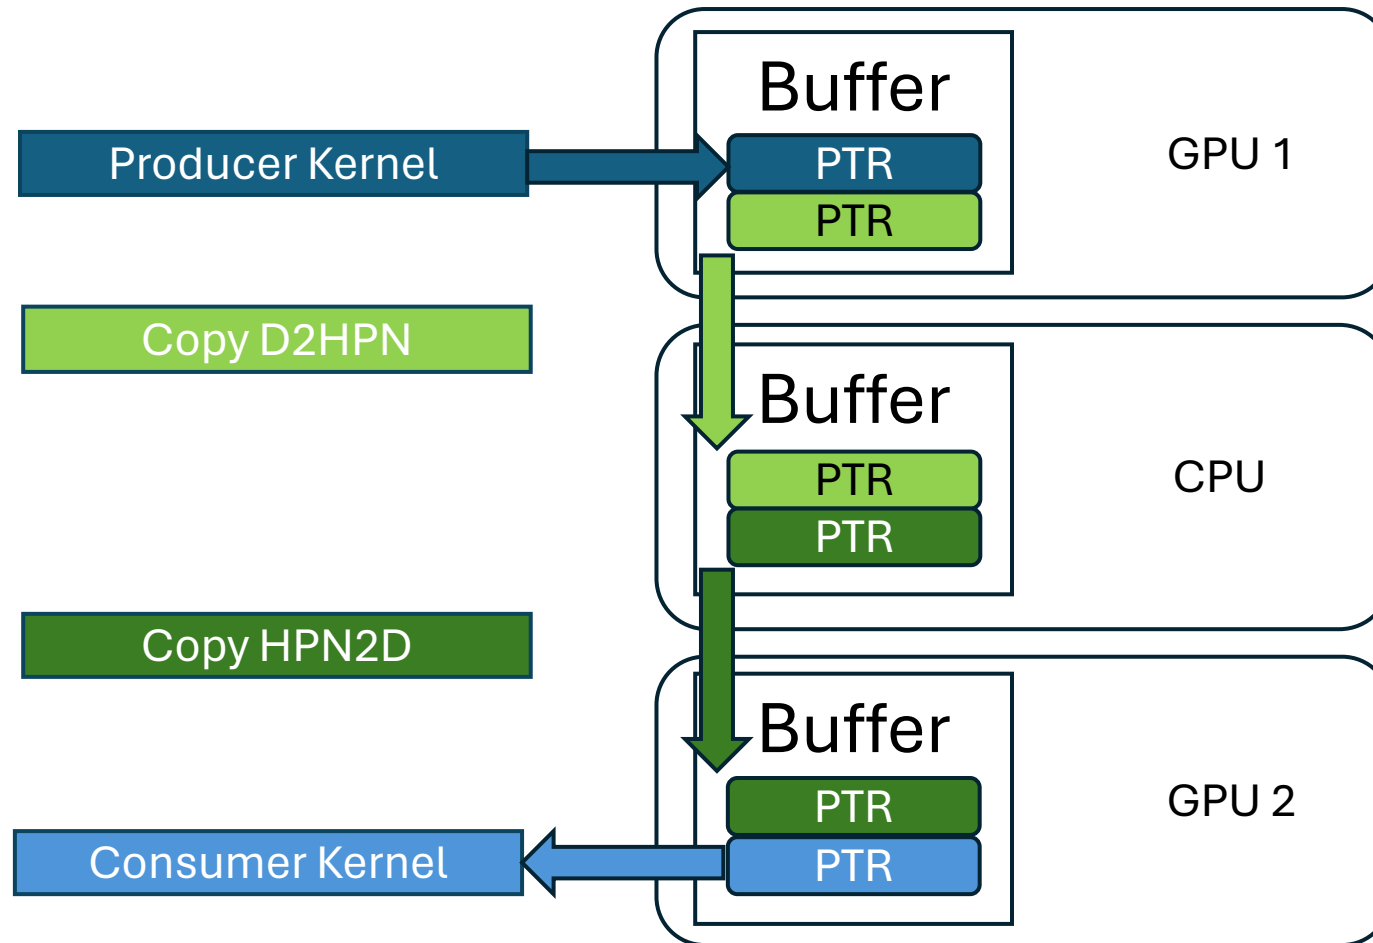
Goal: parallel execution



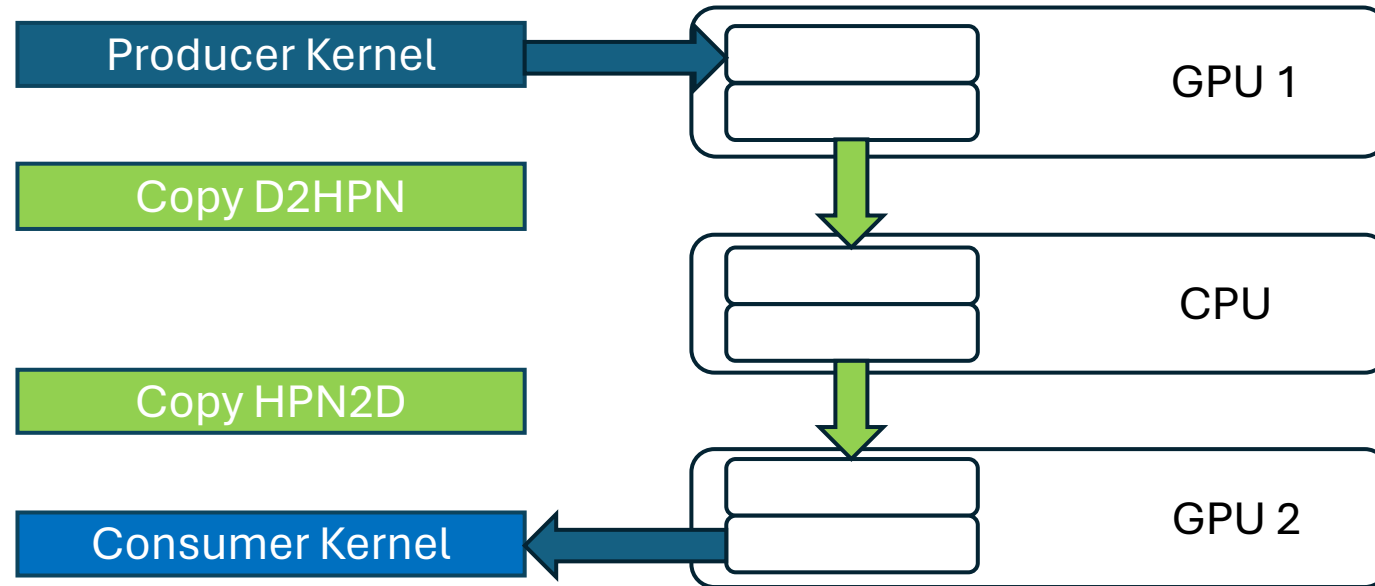
GPU communication utility



GPU communication utility




GPU communication utility



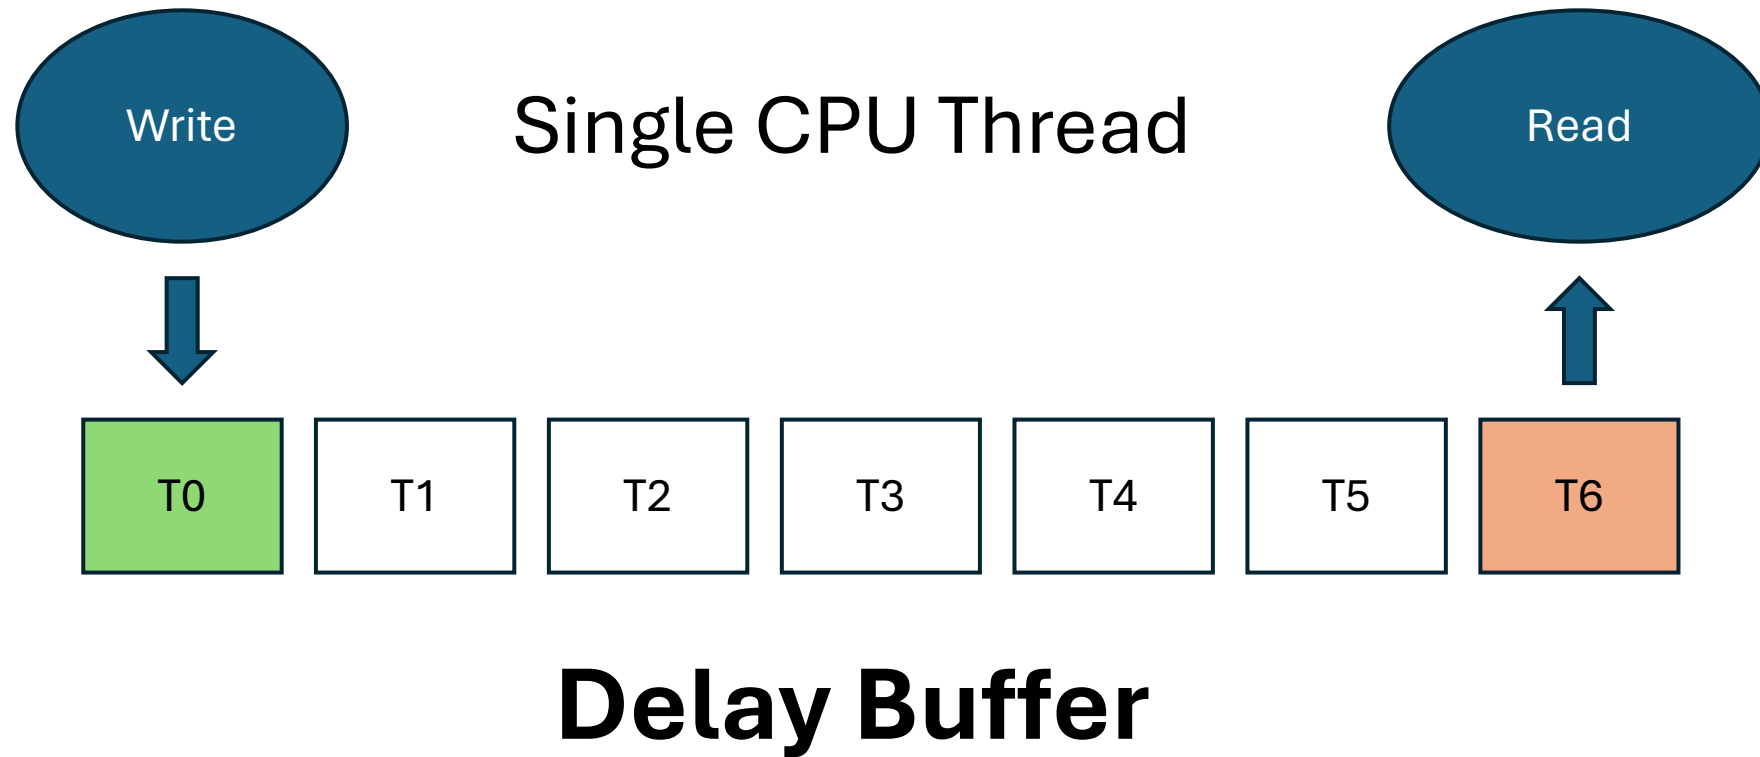
Iteration: 0

END

GPU communication utility

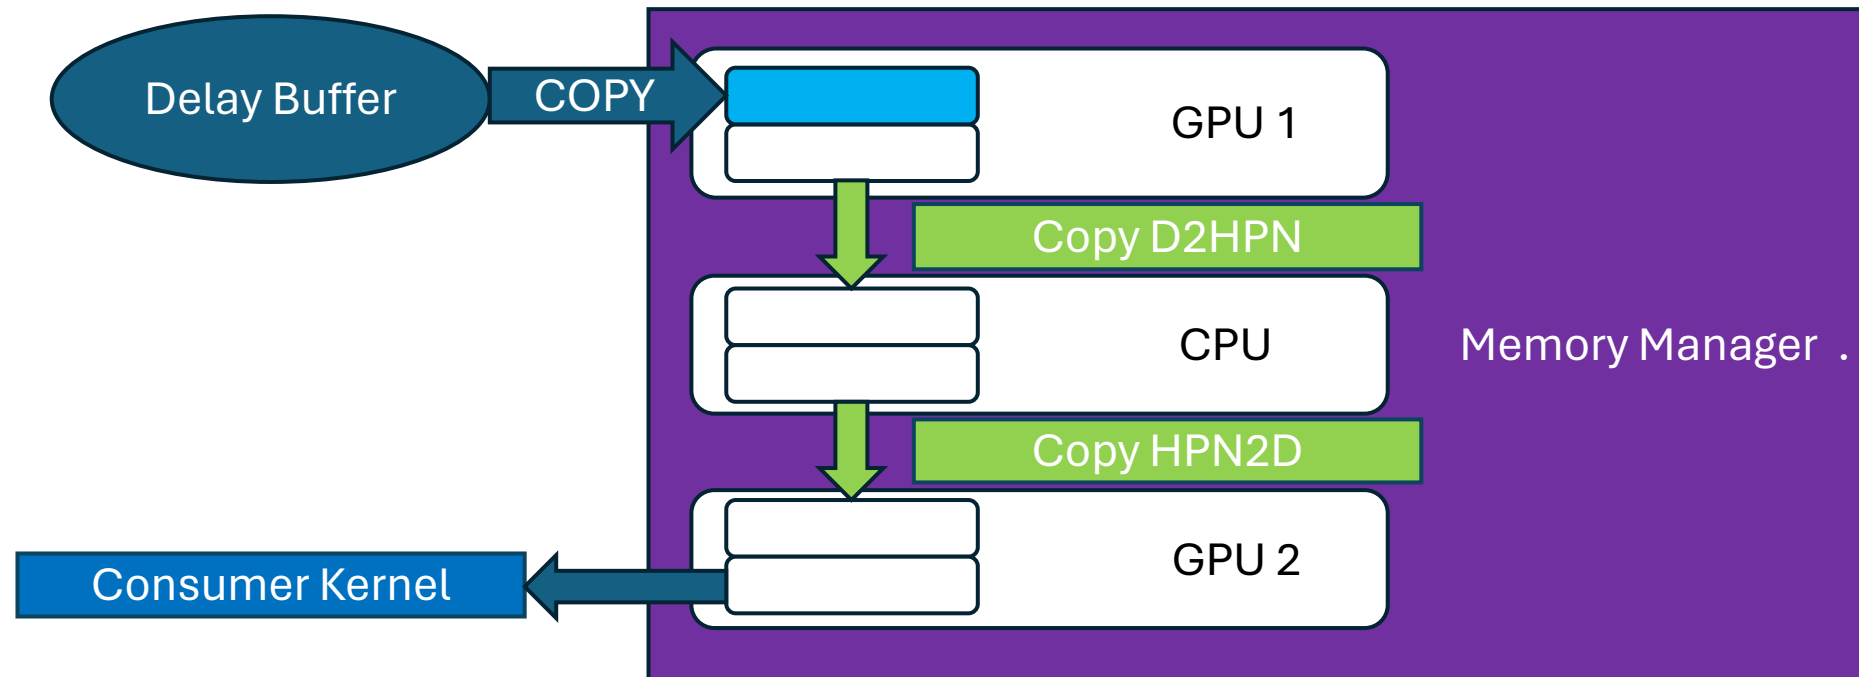


GPU communication utility




GPU communication utility

Memory manager: minimum delay possible

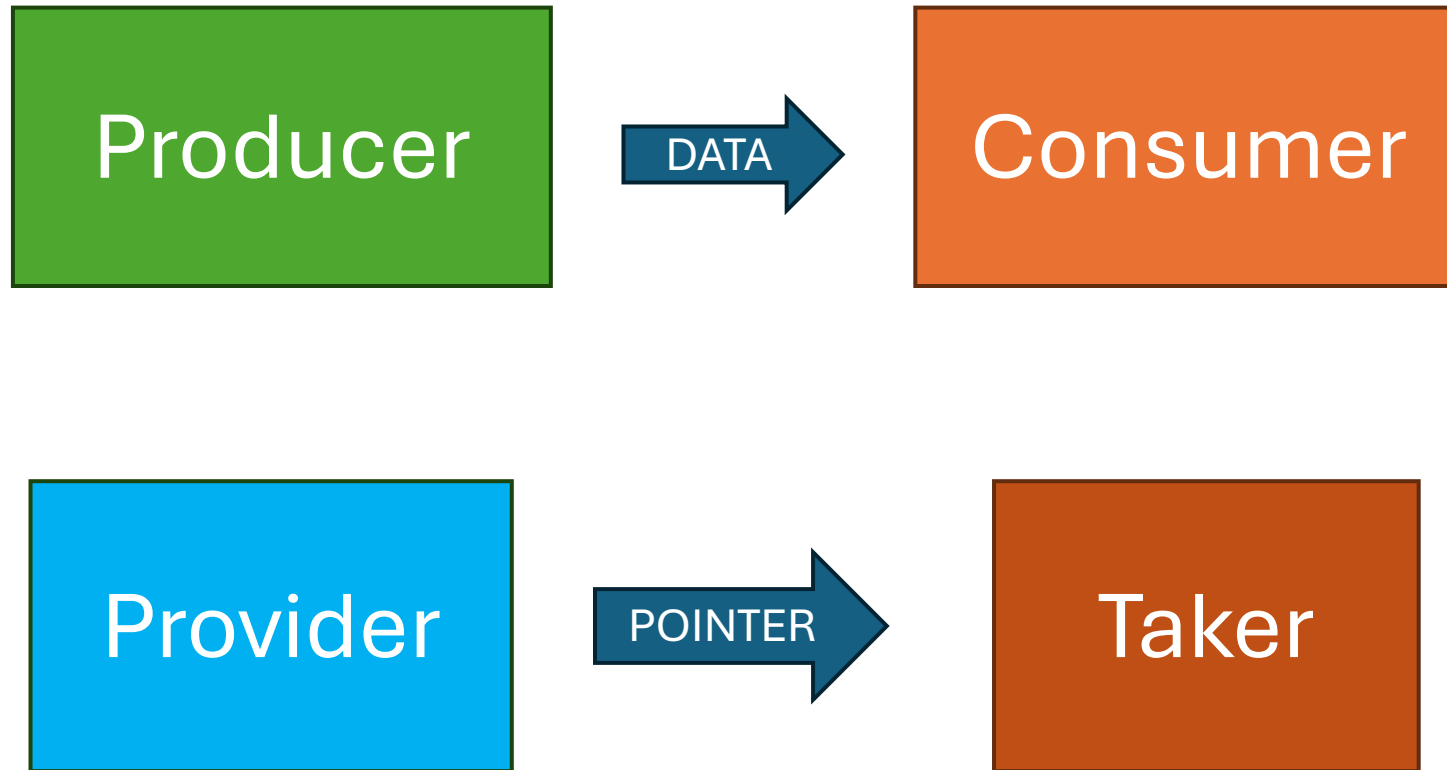


GPU communication utility

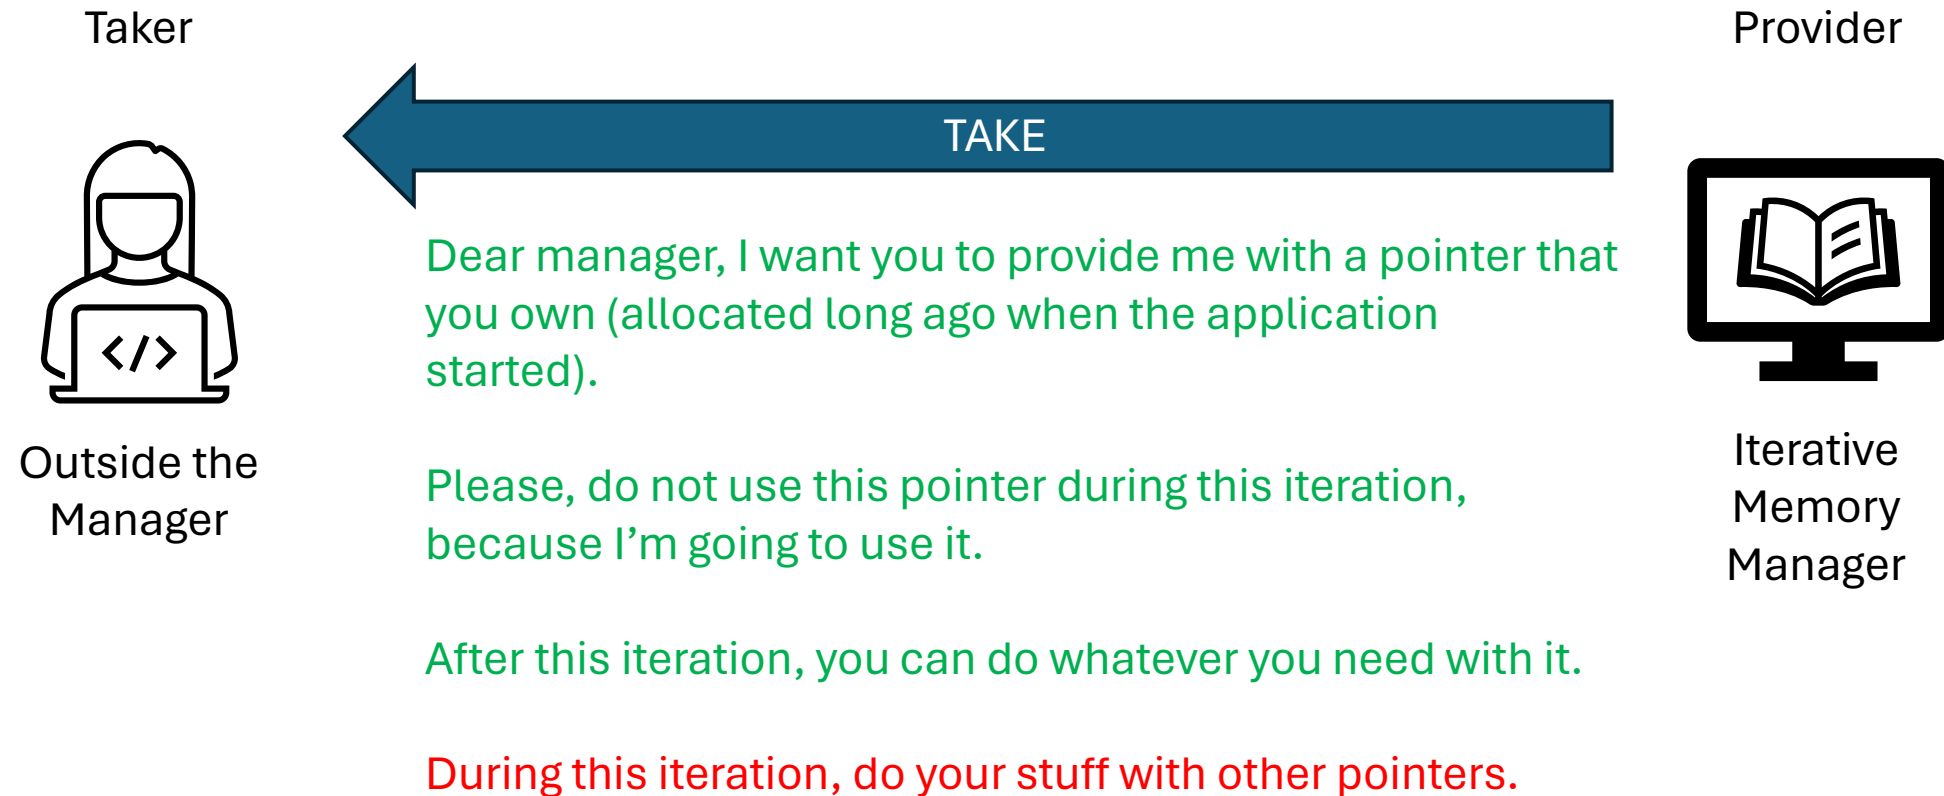
Memory manager: minimum copies and delay possible



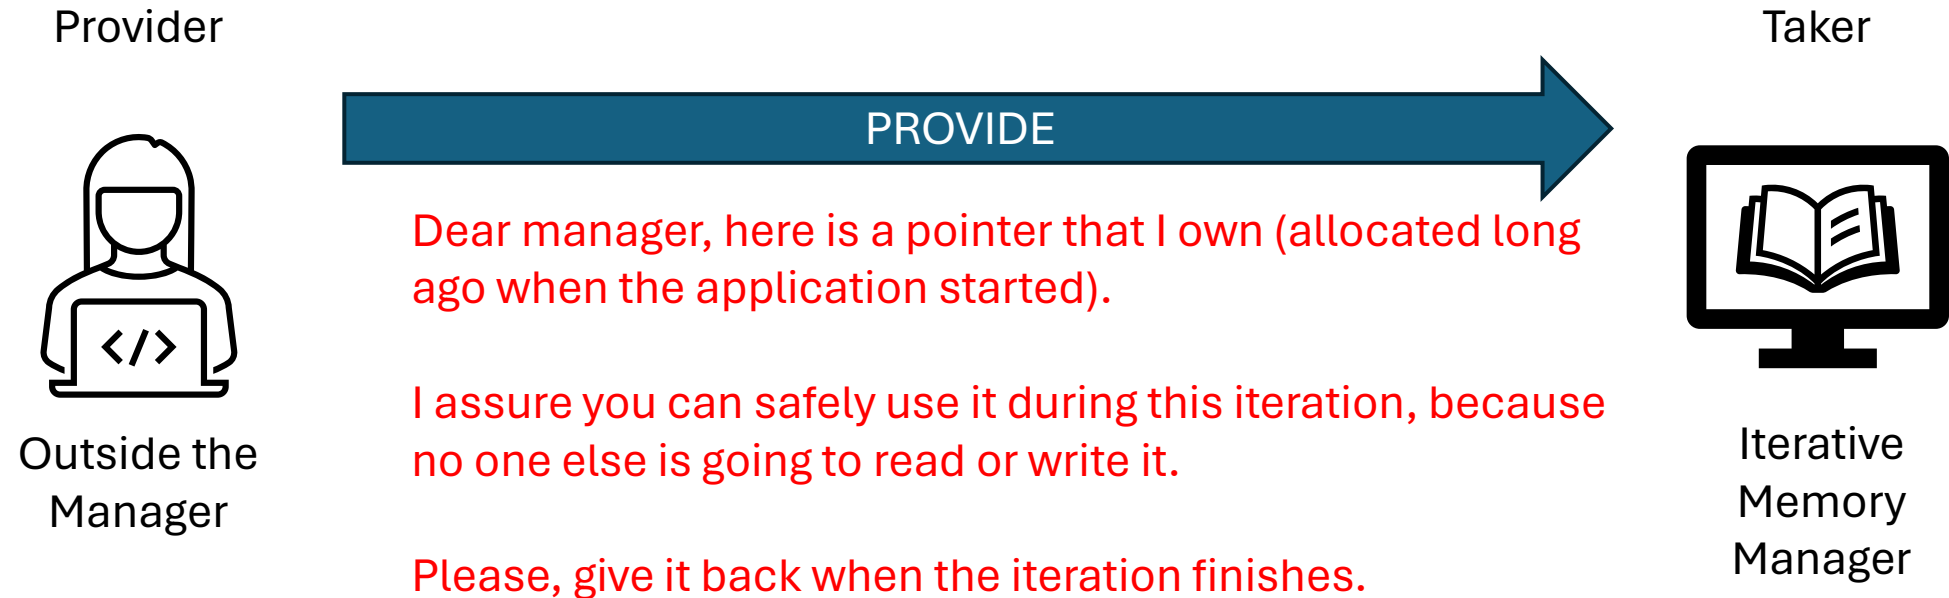
GPU communication utility



GPU communication utility



GPU communication utility



GPU communication utility

| Producer Memory Space | Consumer Memory Space |
|-----------------------|-----------------------|
| ProducerTakes | ConsumerTakes |
| ProducerProvides | ConsumerTakes |
| ProducerTakes | ConsumerProvides |
| ProducerProvides | ConsumerProvides |

GPU communication utility

| Source MS | Destination MS | Flag | Number of copies |
|------------|----------------|---------|------------------------|
| CPU | CPU | HPG2HPG | 0 |
| CPU | CPU Pinned | HPG2HPN | 1 |
| CPU | GPU | HPG2D | 2 CPU -> Pinned -> GPU |
| CPU Pinned | CPU | HPN2HPG | 1 |
| CPU Pinned | CPU Pinned | HPN2HPN | 0 |
| CPU Pinned | GPU | HPN2D | 1 |
| GPU | CPU | D2HPG | 2 |
| GPU | CPU Pinned | D2HPN | 1 |
| GPU | GPU | D2D | 0 |
| GPU | Peer GPU | D2PD | 1 |
| GPU | Fake Peer GPU | D2FPD | 2 GPU -> Pinned -> GPU |



| Producer MS | Consumer MS |
|-------------|-------------|
| Take | Take |
| Provide | Take |
| Take | Provide |
| Provide | Provide |

GPU communication utility

```
struct DataInfo {  
    int numElements;  
    int elemSizeInBytes;  
    MemorySpace memSpace;  
};
```

```
DataInfo producerDataInfo{1024, 4, HostPageable};  
DataInfo consumerDataInfo{1024, 4, Device_1};
```

```
Data ptrToProduce(producerDataInfo);  
Data ptrToConsume(consumerDataInfo);
```

```
enum Actions { ProducerProvides, ProducerTakes, ConsumerProvides, ConsumerTakes};
```


GPU communication utility

//Initialization

```
MemoryManager<ProducerProvides, ConsumerProvides> manager(producerDataInfo, consumerDataInfo);
```

```
int delay = manager.getTotalDelay(); // Query delay generated by the manager
```

```
manager.manage(ptrToProduce, ptrToConsume); // Usage
```

GPU communication utility

```
// Initialization
```

```
MemoryManager<ProducerTakes, ConsumerTakes> manager(producerDataInfo, consumerDataInfo);
```

```
auto [ptrToProduce, ptrToConsume] = manager.manage(); // Usage
```

GPU communication utility

// Initialization

```
MemoryManager<ProducerTakes, ConsumerProvides> manager(producerDataInfo, consumerDataInfo);
```

```
ptrToProduce = manager.manage(ptrToConsume); // Usage
```

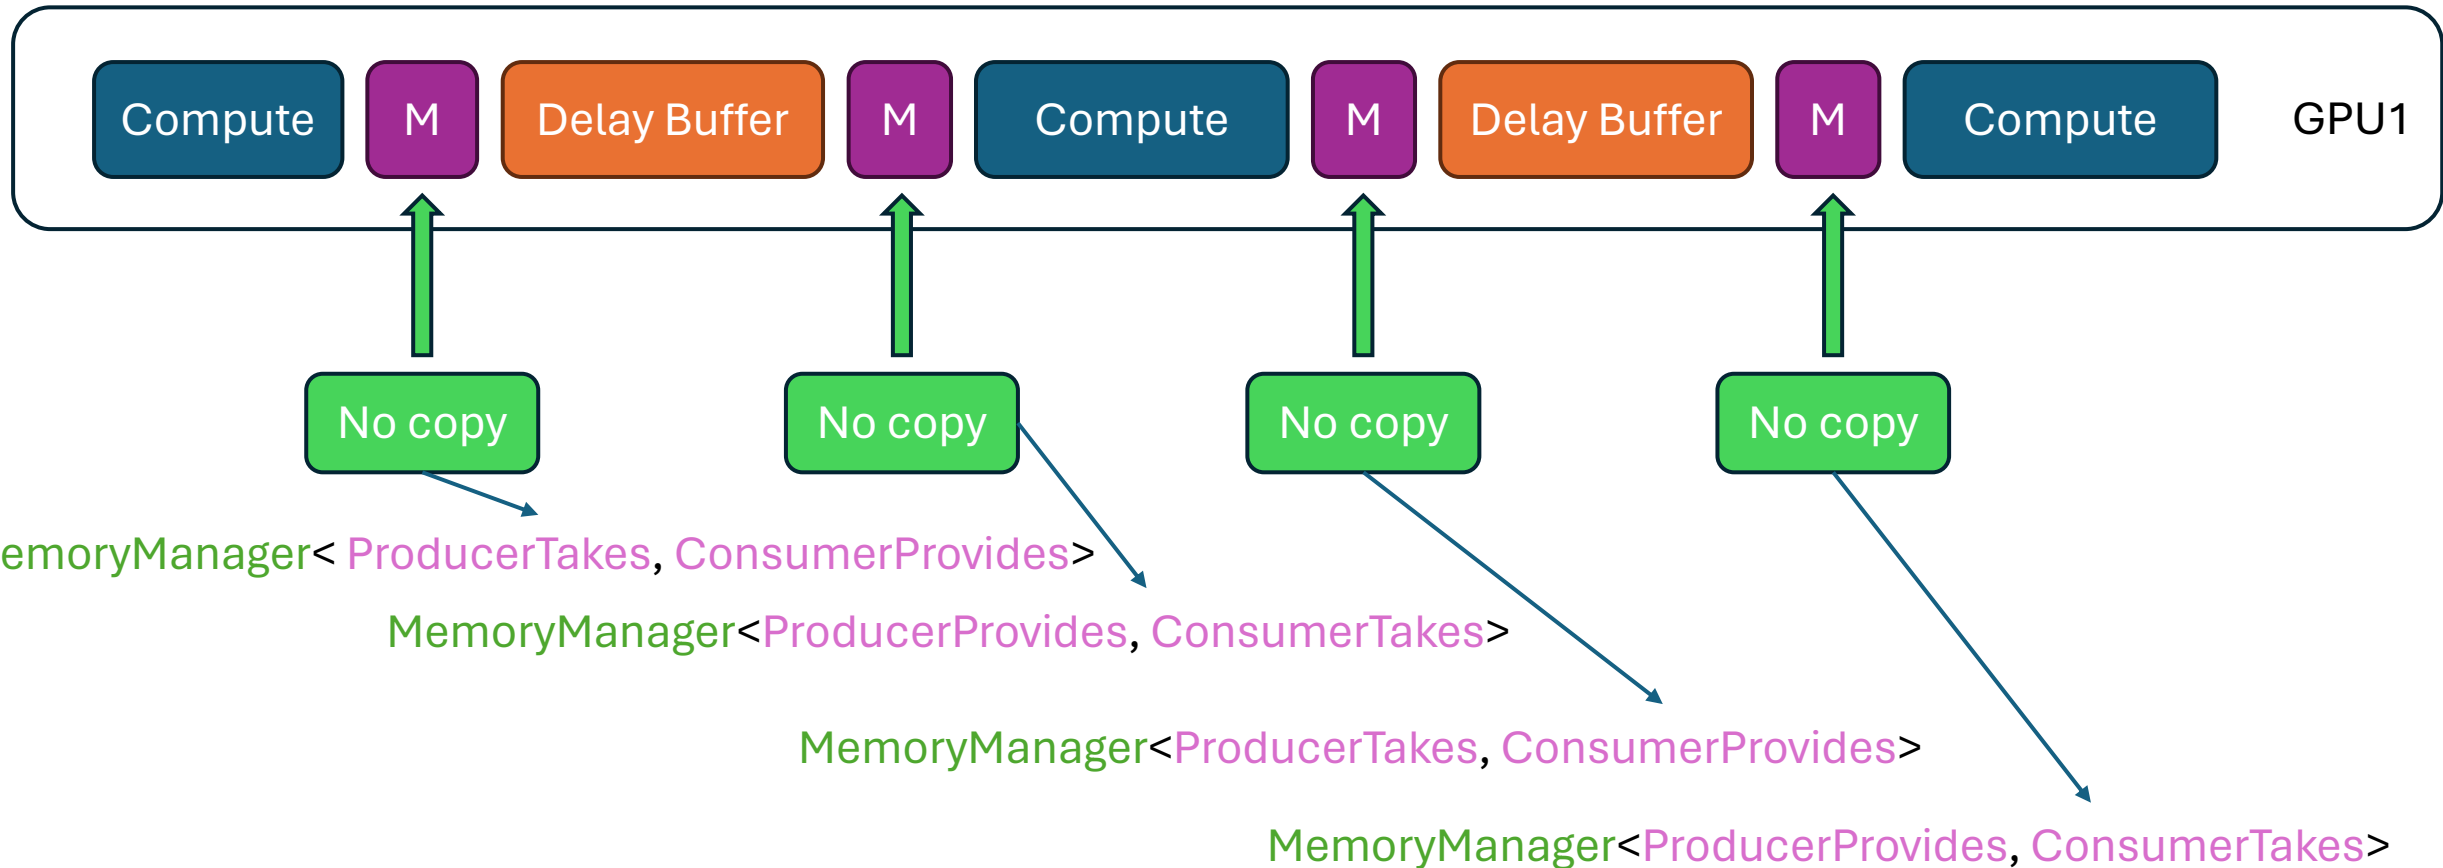
GPU communication utility

```
// Initialization
```

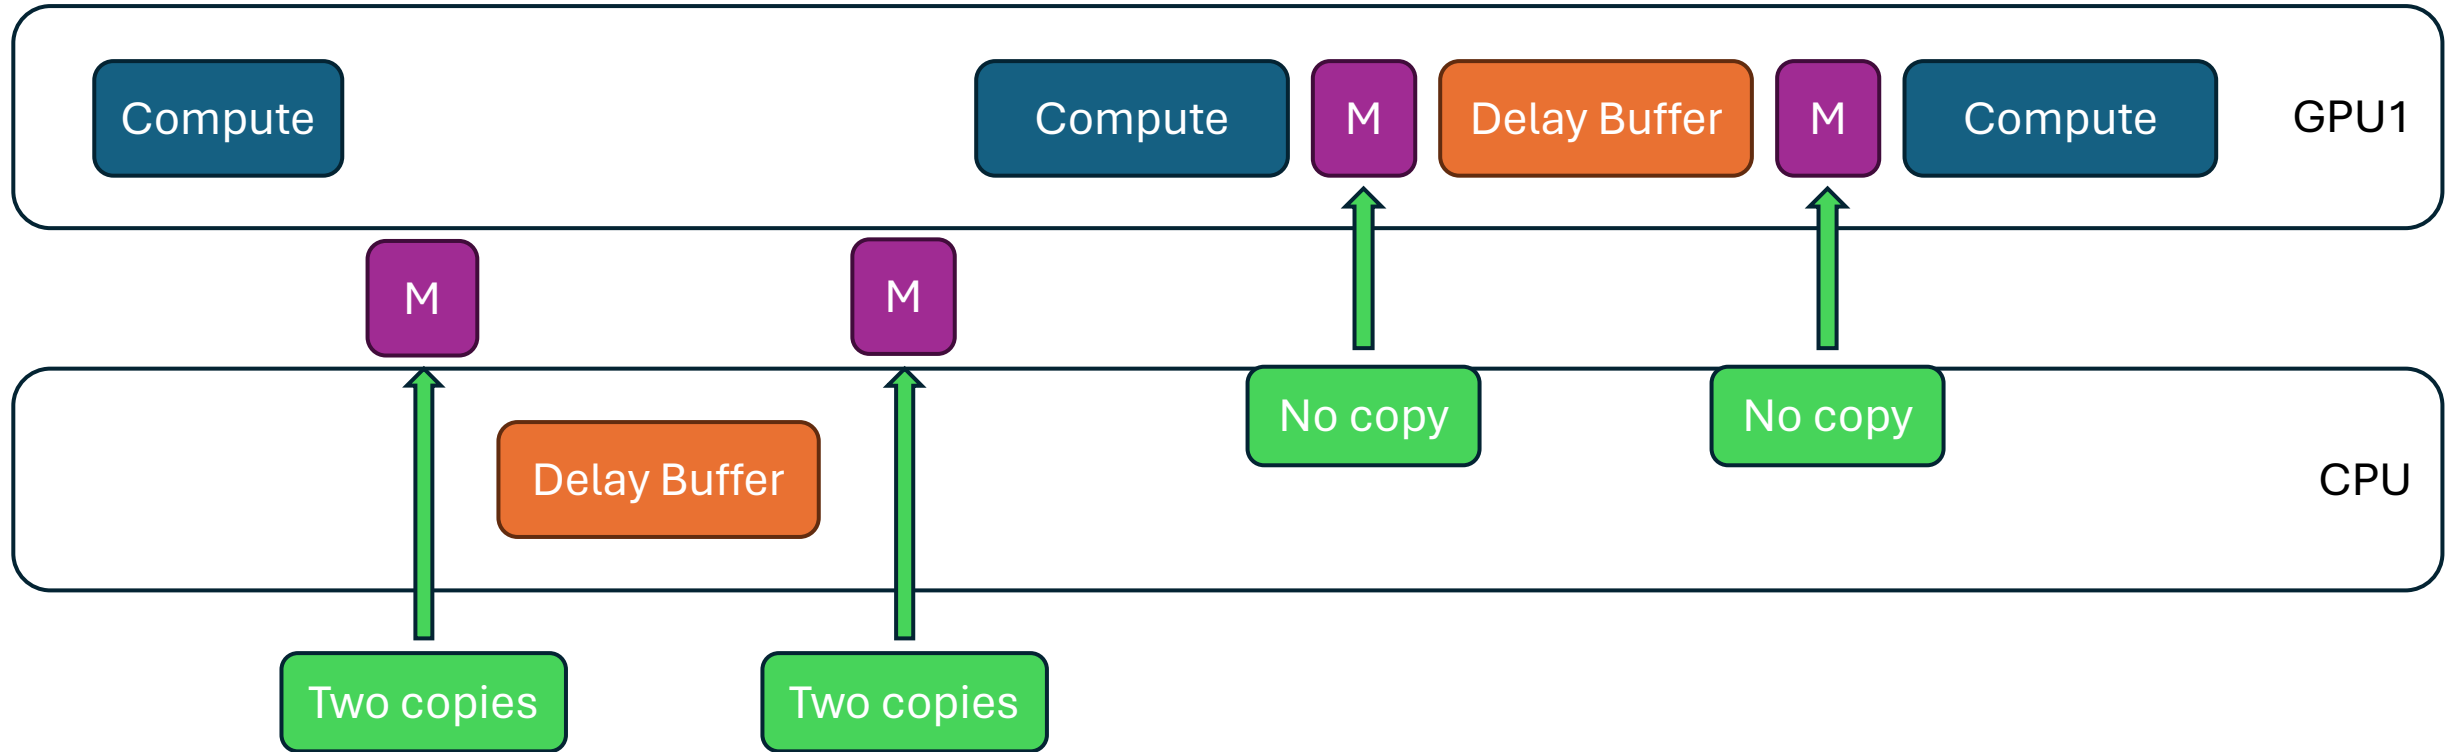
```
MemoryManager<ProducerProvides, ProducerTakes> manager(producerDataInfo, consumerDataInfo);
```

```
ptrToConsume = manager.manage(ptrToProduce); // Usage
```

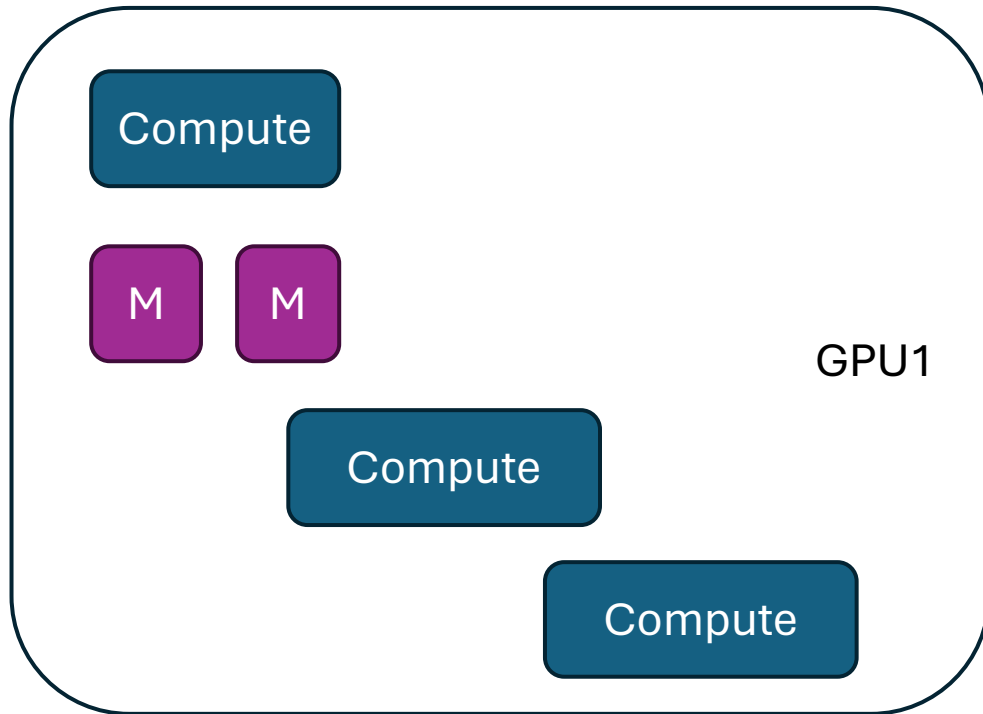
GPU communication utility



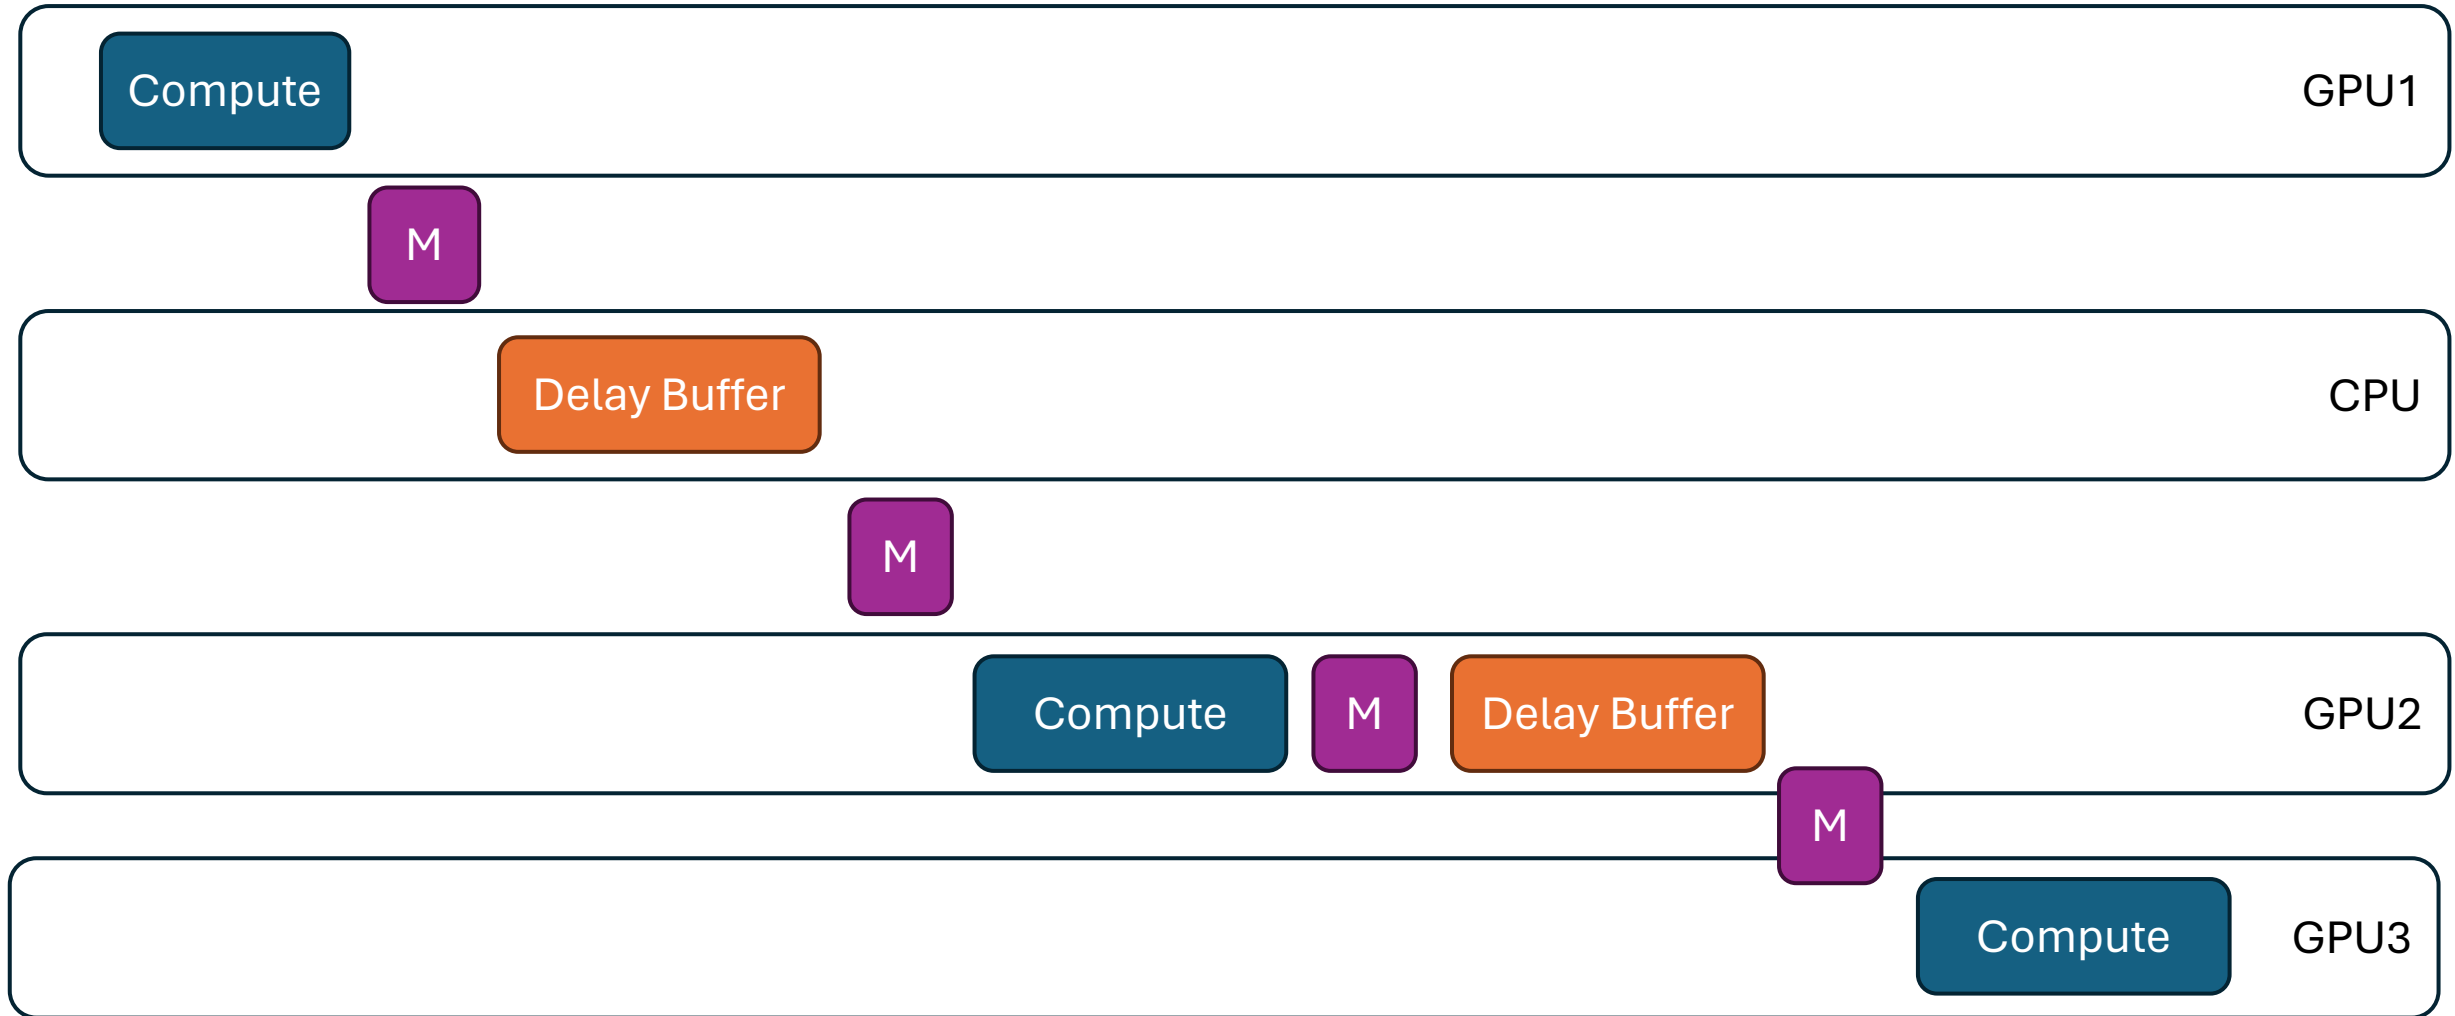
GPU communication utility



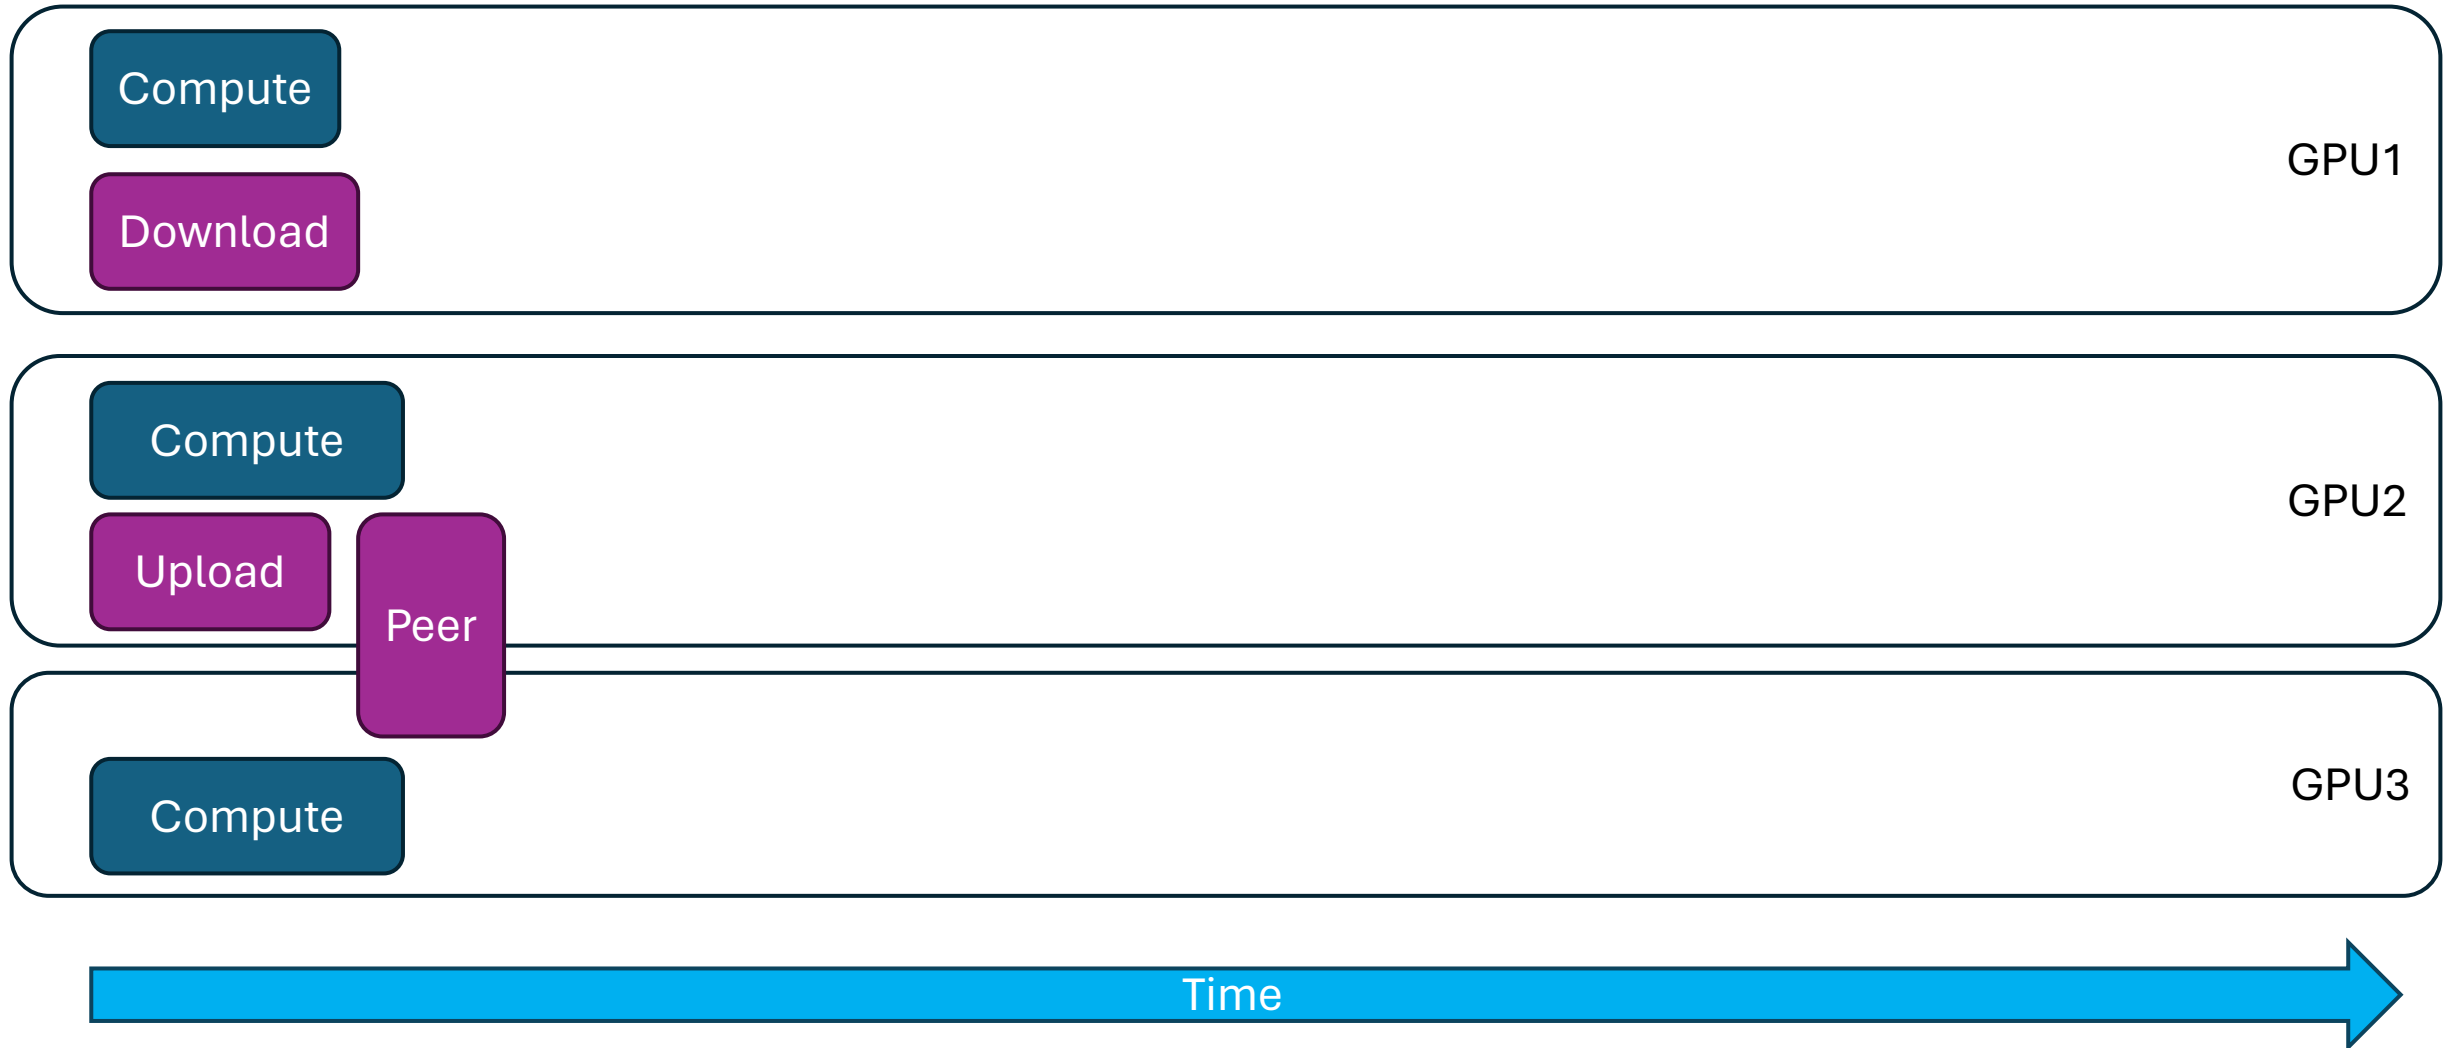
GPU communication utility





GPU communication utility



GPU communication utility



GPU communication utility



GPU communication utility

Libraries that provide solutions to similar problems

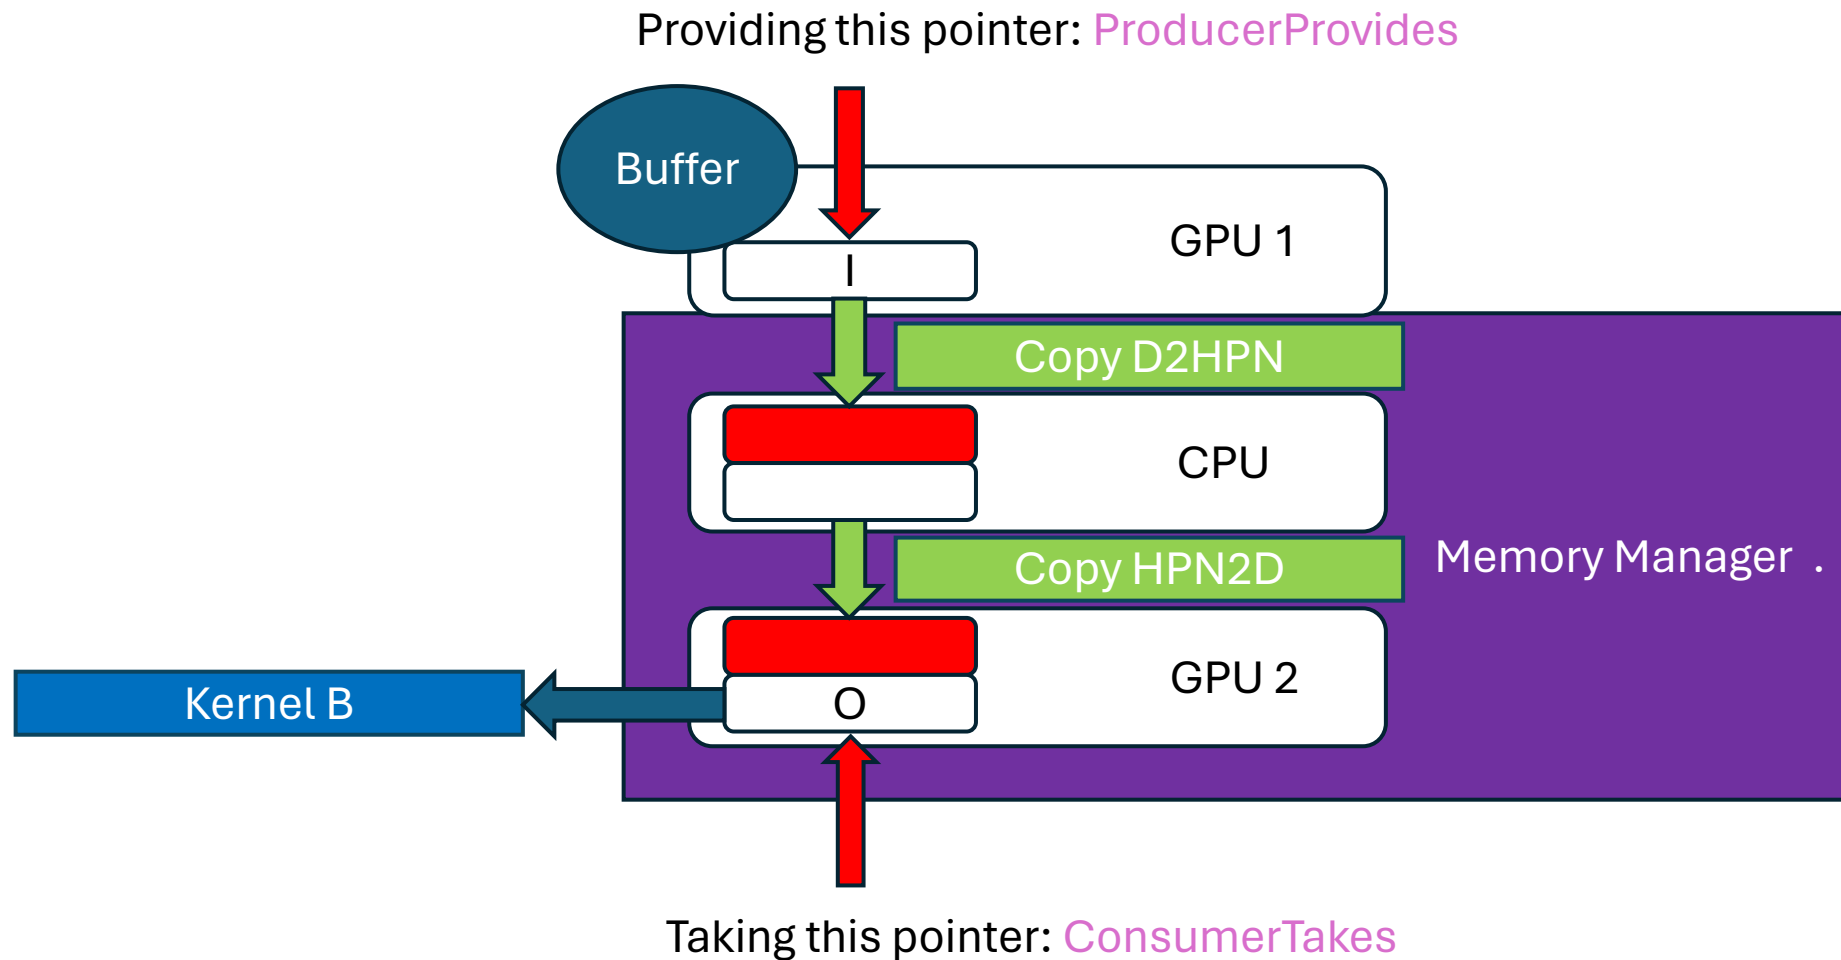
Image processing graphs:

[DeepStream SDK | NVIDIA Developer | NVIDIA Developer](#)

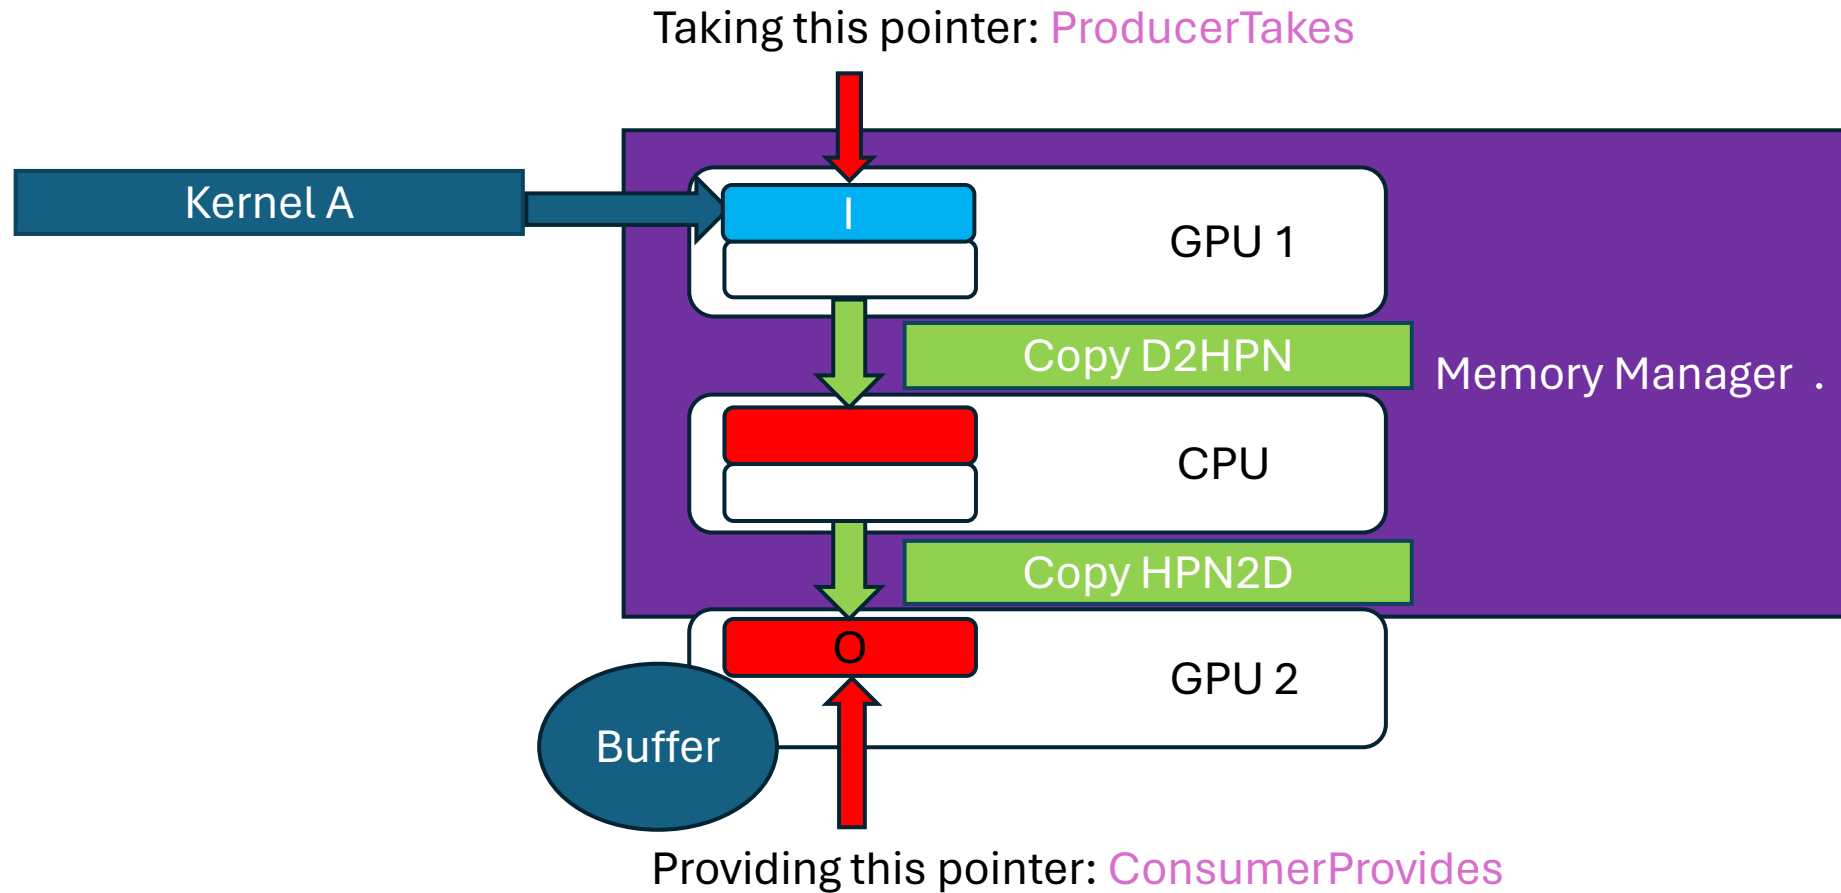
Multi-gpu and multi-node CUDA programs (training DNNs and more):

[NVIDIA Collective Communications Library \(NCCL\) | NVIDIA Developer](#)

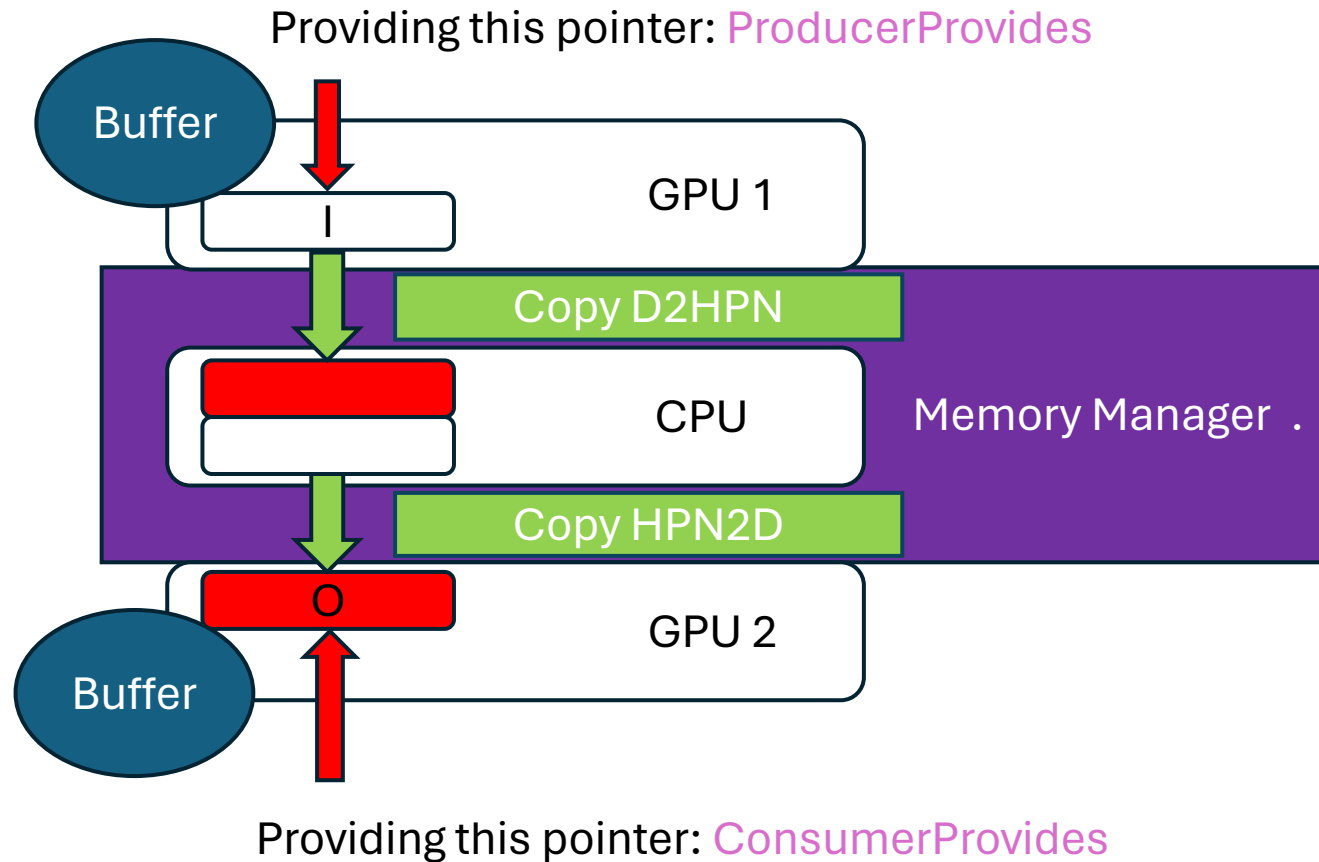
GPU communication utility



GPU communication utility

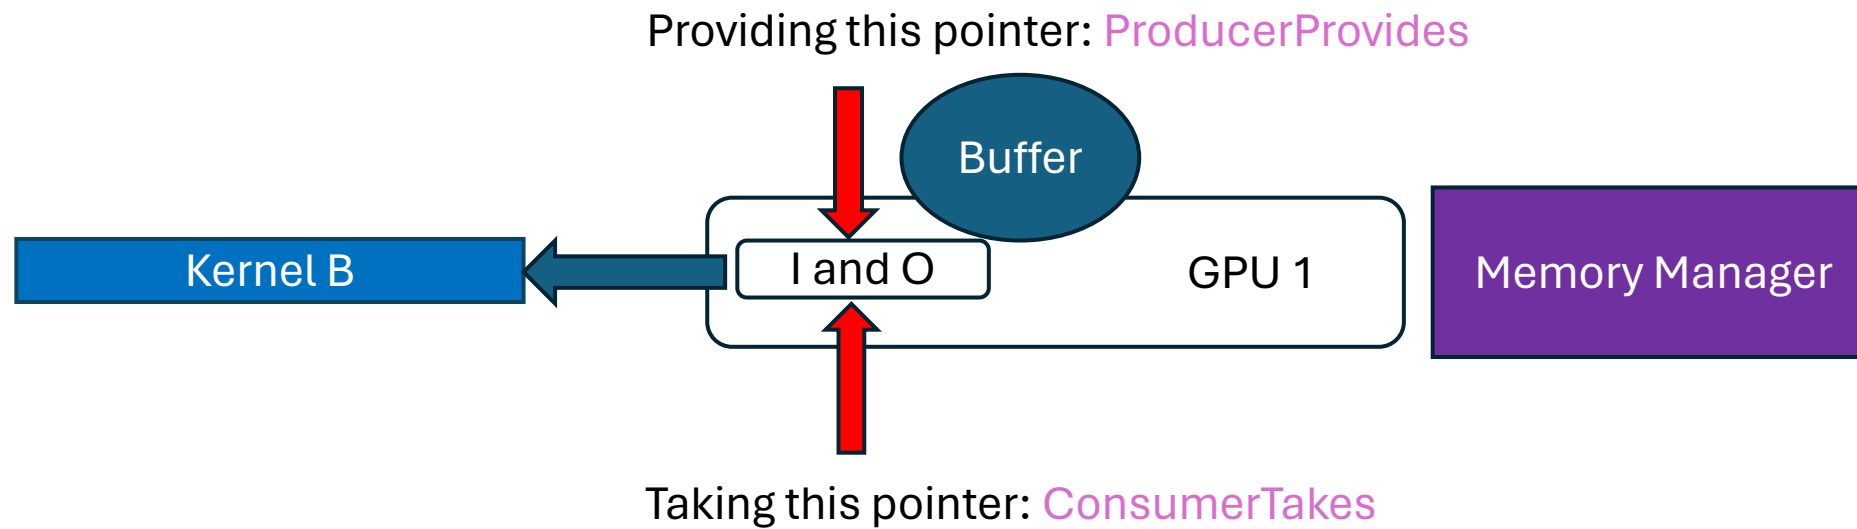


GPU communication utility



GPU communication utility

Zero copy case



GPU communication utility

Rules to know when to use **Provide**:

- If the Producer code it's a Delay Buffer.
- Equivalent to saying: if the Producer code **does not modify** the data during this iteration.
- Then Producer **must** Provide: **ProducerProvides**
- If the Consumer code it's a Delay Buffer.
- Equivalent to saying: if the Consumer code **does not modify** the data during this iteration.
- Then Consumer **must** Provide: **ConsumerProvides**

GPU communication utility

Rules to know when to use **Take**:

- If the Producer code it's a Kernel or asynchronous CPU code.
- Equivalent to saying: if the Producer code **does modify** the data during this iteration.
- Then Producer **must** Take: **ProducerTakes**
- If the Consumer code it's a Kernel or asynchronous CPU code.
- Equivalent to saying: if the Consumer code **does modify** the data during this iteration.
- Then Consumer **must** Take: **ConsumerTakes**

Lecture overview

Summary of already
seen concepts

Main story: libraries

Use cases:

- GPU communication utility
- GPU kernel libraries







GPU kernel libraries

OpenSource personal project:

[morousg/cvGPUSpeedup](#): A faster implementation of OpenCV-CUDA that uses OpenCV objects, and more! ([github.com](#))

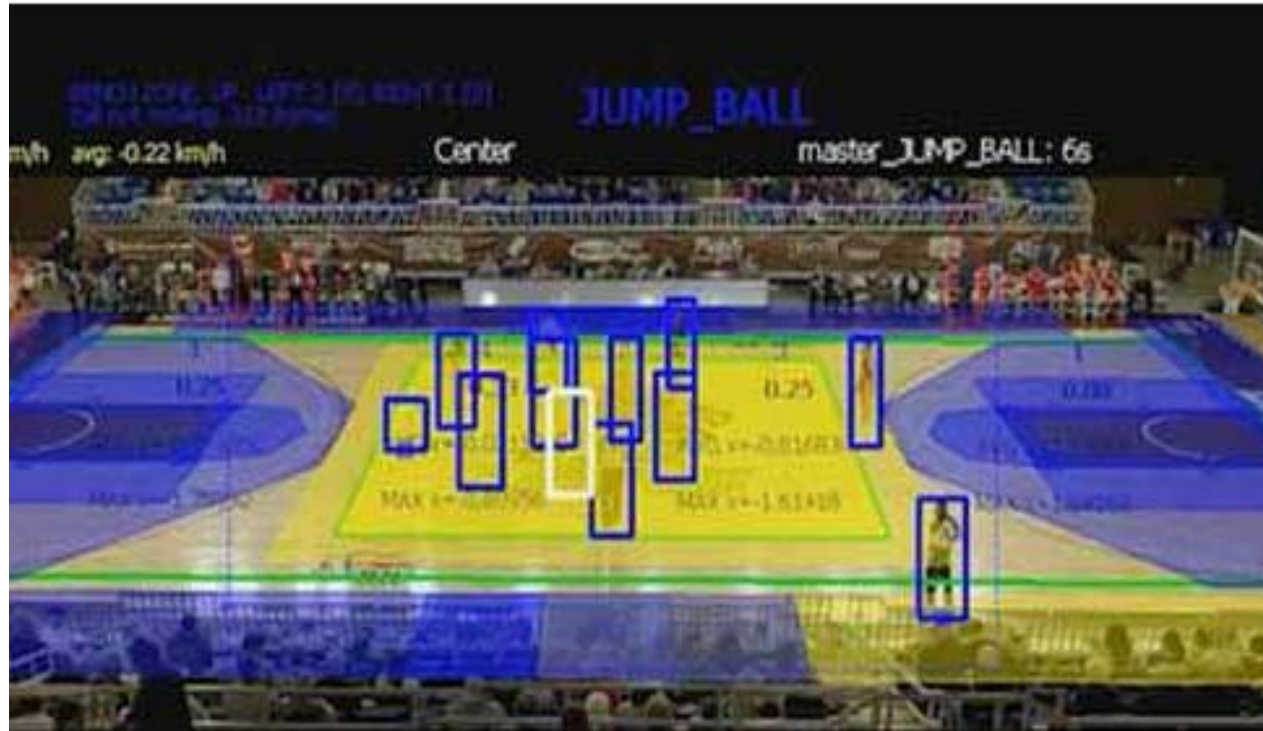
GPU kernel libraries

- The inference image preprocessing
- The image processing pipeline

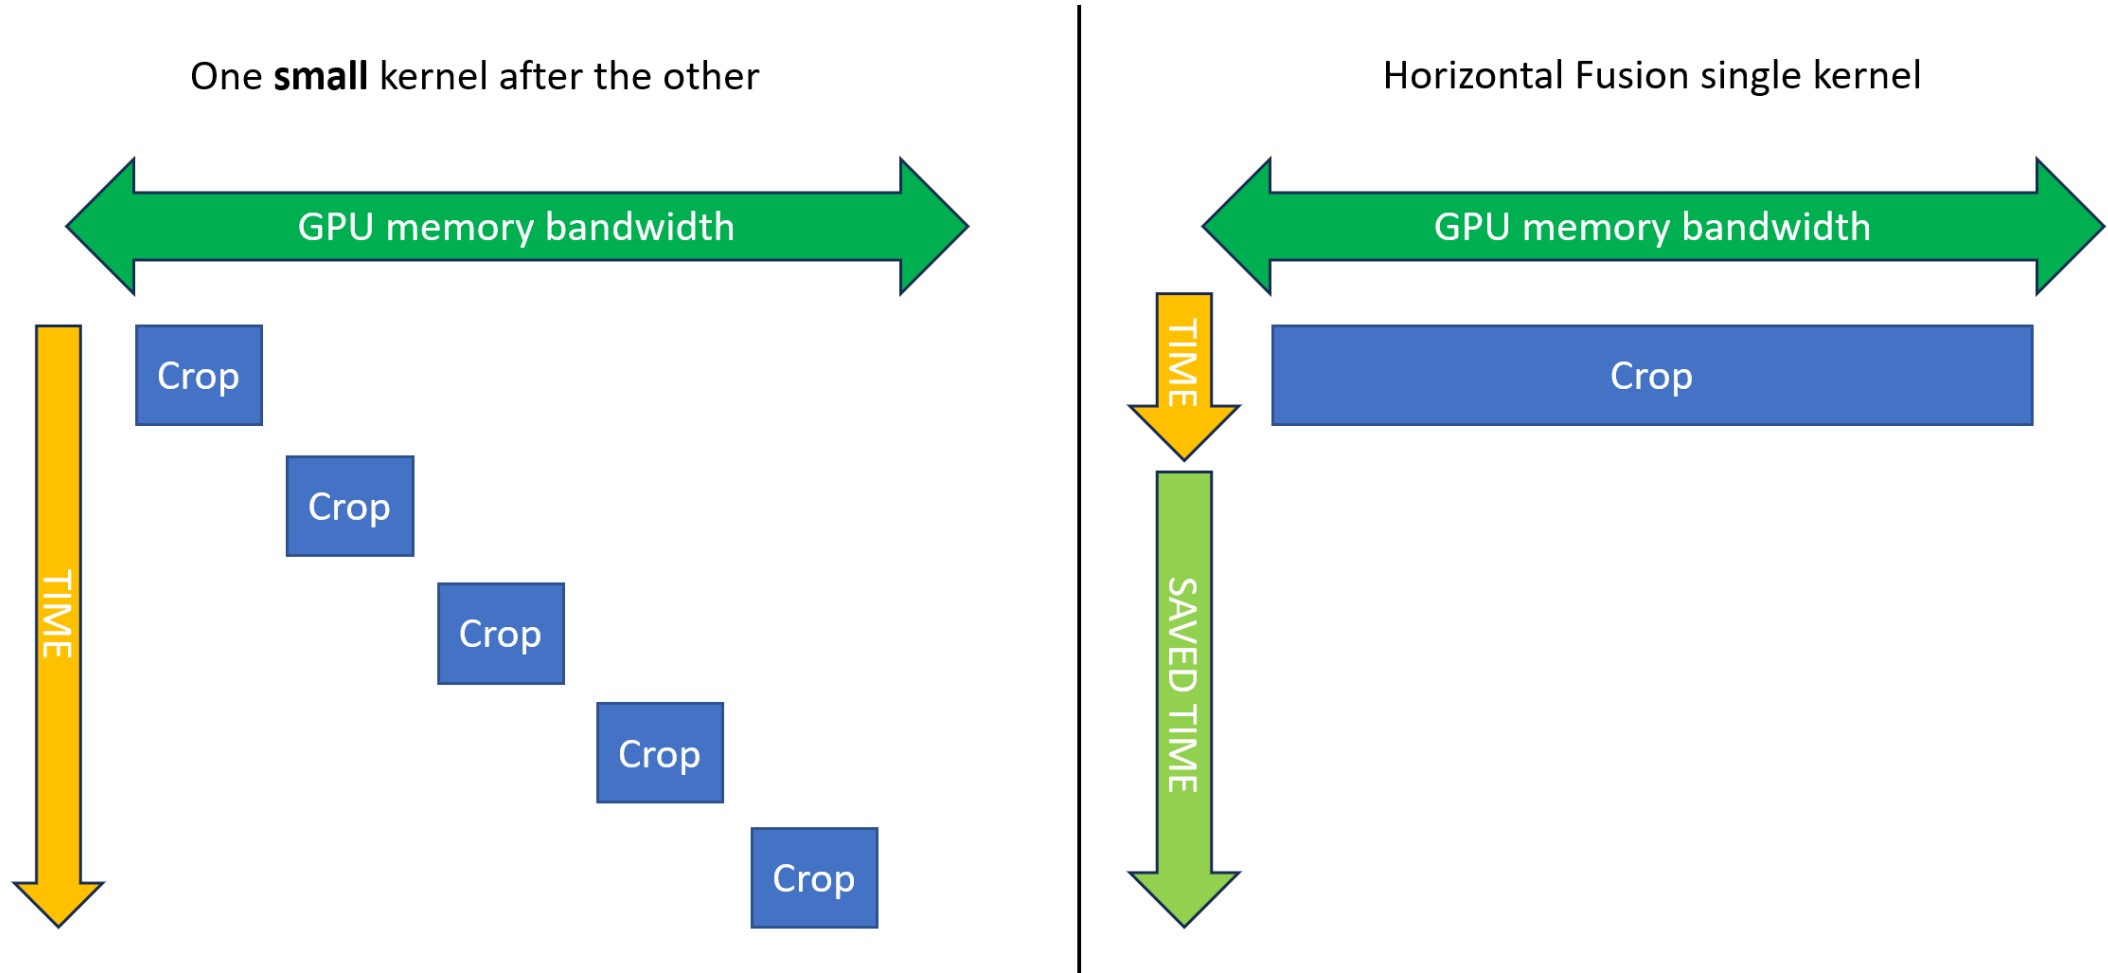


GPU kernel libraries:

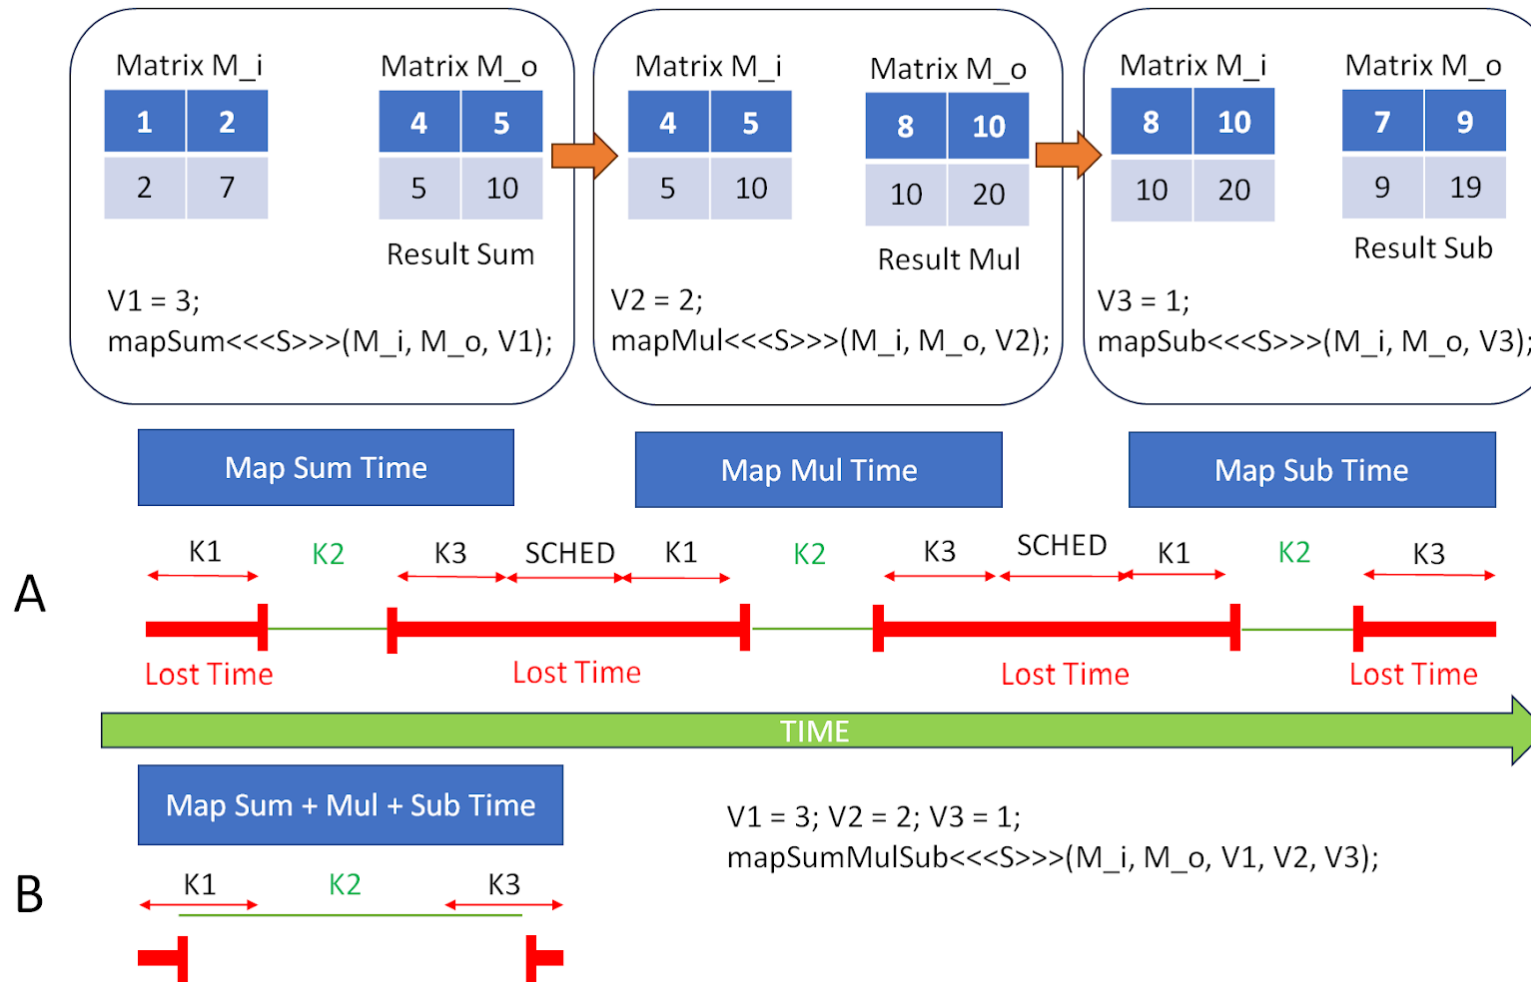
- AI team uses OpenCV-CUDA -> extremely memory bound kernels
- Many small independent and dependent kernels



GPU kernel libraries: Horizontal Fusion



GPU kernel libraries: Vertical Fusion



GPU kernel libraries: Fusion



Vertical Fusion for
Data Dependent
Kernels



Horizontal Fusion for
Data Independent
Kernels

GPU kernel libraries: novel concepts implemented in cvGPUSpeedup library

Generic Vertical Fusion (GVF)

Divergent Horizontal Fusion (DHF)

Backwards Generic Vertical Fusion (BGVF)

Automatic Thread Coarsening (ATC)

GPU kernel libraries

- **Typical approach for Vertical Fusion:**
 - Library defines a set of Vertically Fused Kernels (`__global__` functions)
 - Explicitly chosen by the user
 - Implicitly chosen by a runtime
- **Our proposal:**
 - The library provides a set of fusionable `__device__` functions
 - The final user, defines the chain of functions that go into the kernel, **without knowing it.**
 - We call this **Generic Vertical Fusion**



GPU kernel libraries: Generic Vertical Fusion

No Fusion




// OpenCV version

```
cv::cuda::resize(d_input(crop), d_up, targetRes, 0., 0., cv::INTER_LINEAR, cv_stream);  
d_up.convertTo(d_temp, CV_32FC3, alpha, cv_stream);  
cv::cuda::subtract(d_temp, val_sub, d_temp2, cv::noArray(), -1, cv_stream);  
cv::cuda::divide(d_temp2, val_div, d_temp, 1.0, -1, cv_stream);  
cv::cuda::split(d_temp, d_output, cv_stream);
```

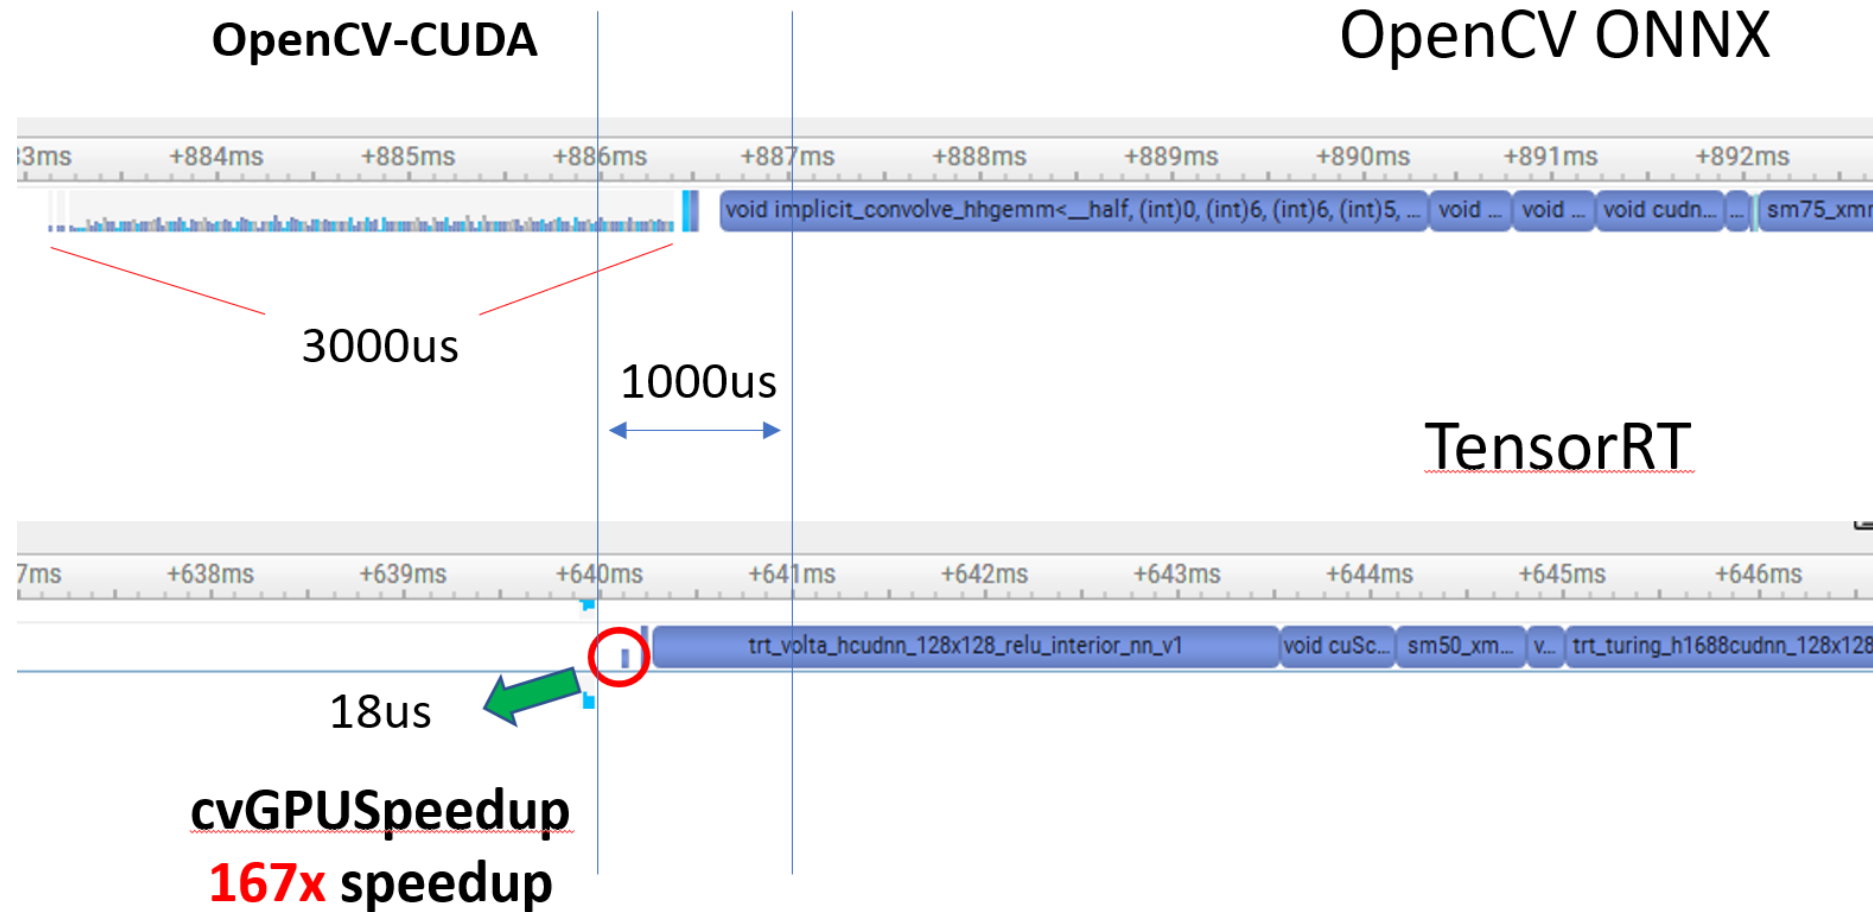
// cvGPUSpeedup version

```
cv::Scalar val_alpha(alpha, alpha, alpha);  
cvGS::executeOperations(cv_stream,  
    cvGS::resize<CV_8UC3, cv::INTER_LINEAR>(d_input(crop), targetRes, 0., 0.),  
    cvGS::convertTo<CV_8UC3, CV_32FC3>(),  
    cvGS::multiply<CV_32FC3>(val_alpha),  
    cvGS::subtract<CV_32FC3>(val_sub),  
    cvGS::divide<CV_32FC3>(val_div),  
    cvGS::split<CV_32FC3>(d_output));
```

GVF

- Same variable 
- Memory savings 
- Different params 

GPU kernel libraries: Cr,Res,Nor for 50 Crops



GPU kernel libraries: CircularTensor

```
cvGS::CircularTensor<InputCVType, CircularTensorCVType, NUM_CHANNELS, BATCH,  
                    cvGS::CircularTensorOrder::NewestFirst> myTensor(WIDTH, HEIGHT);
```

```
myTensor.update(cv_stream,  
               cvGS::resize<...>(newImage, ...),  
               cvGS::convertTo<...>(...),  
               cvGS::split<...>(myTensor.ptr().data)); // We may look for a way to avoid this
```

// Now you can send the raw data to inference

```
network.forward(myTensor.ptr().data, cv_stream);
```

New image



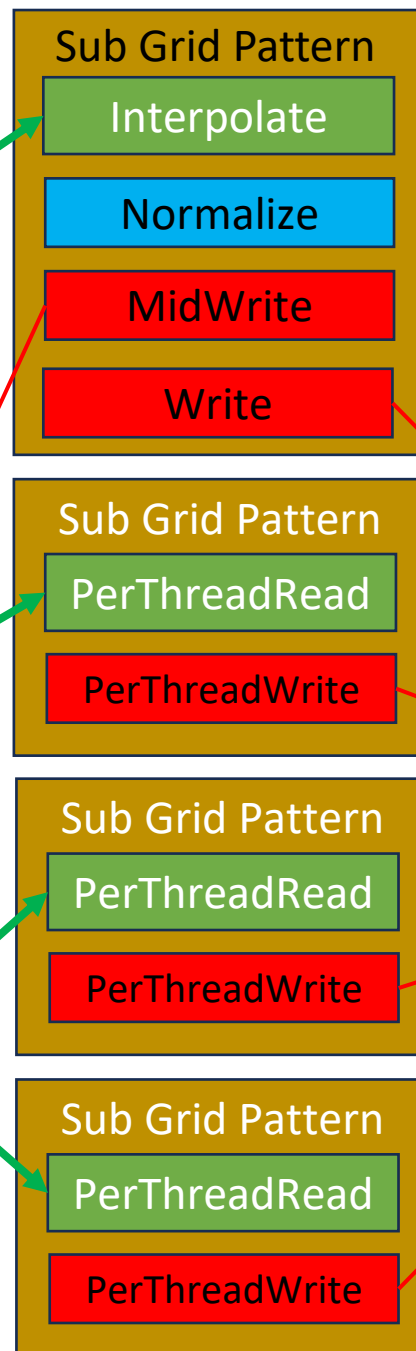

CircularTensor

Using: GVF + DHF

Fused Kernel


Newest

Oldest

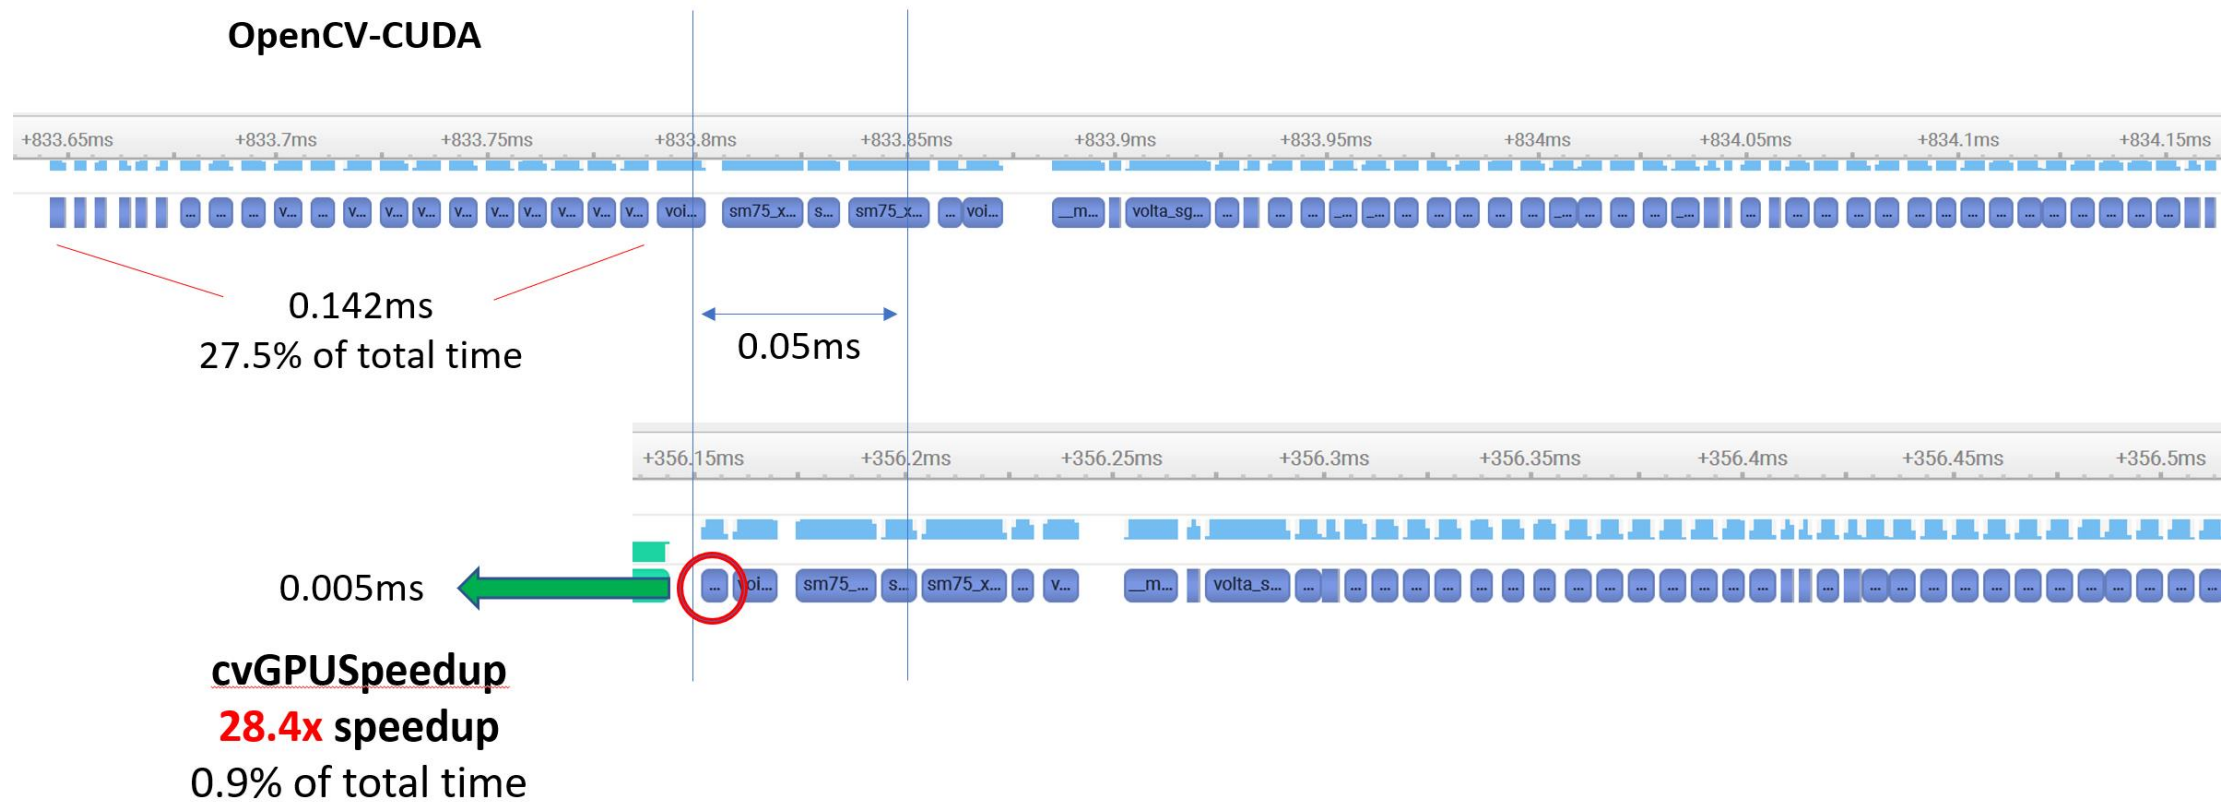


Newest

Oldest




GPU kernel libraries



GPU kernel libraries: image processing

- Hand made kernels
- Mostly memory bound -> repetitive optimizations
- Read uchar4, process float4, store uchar4 -> loss of image quality



GPU kernel libraries: graph total fusion

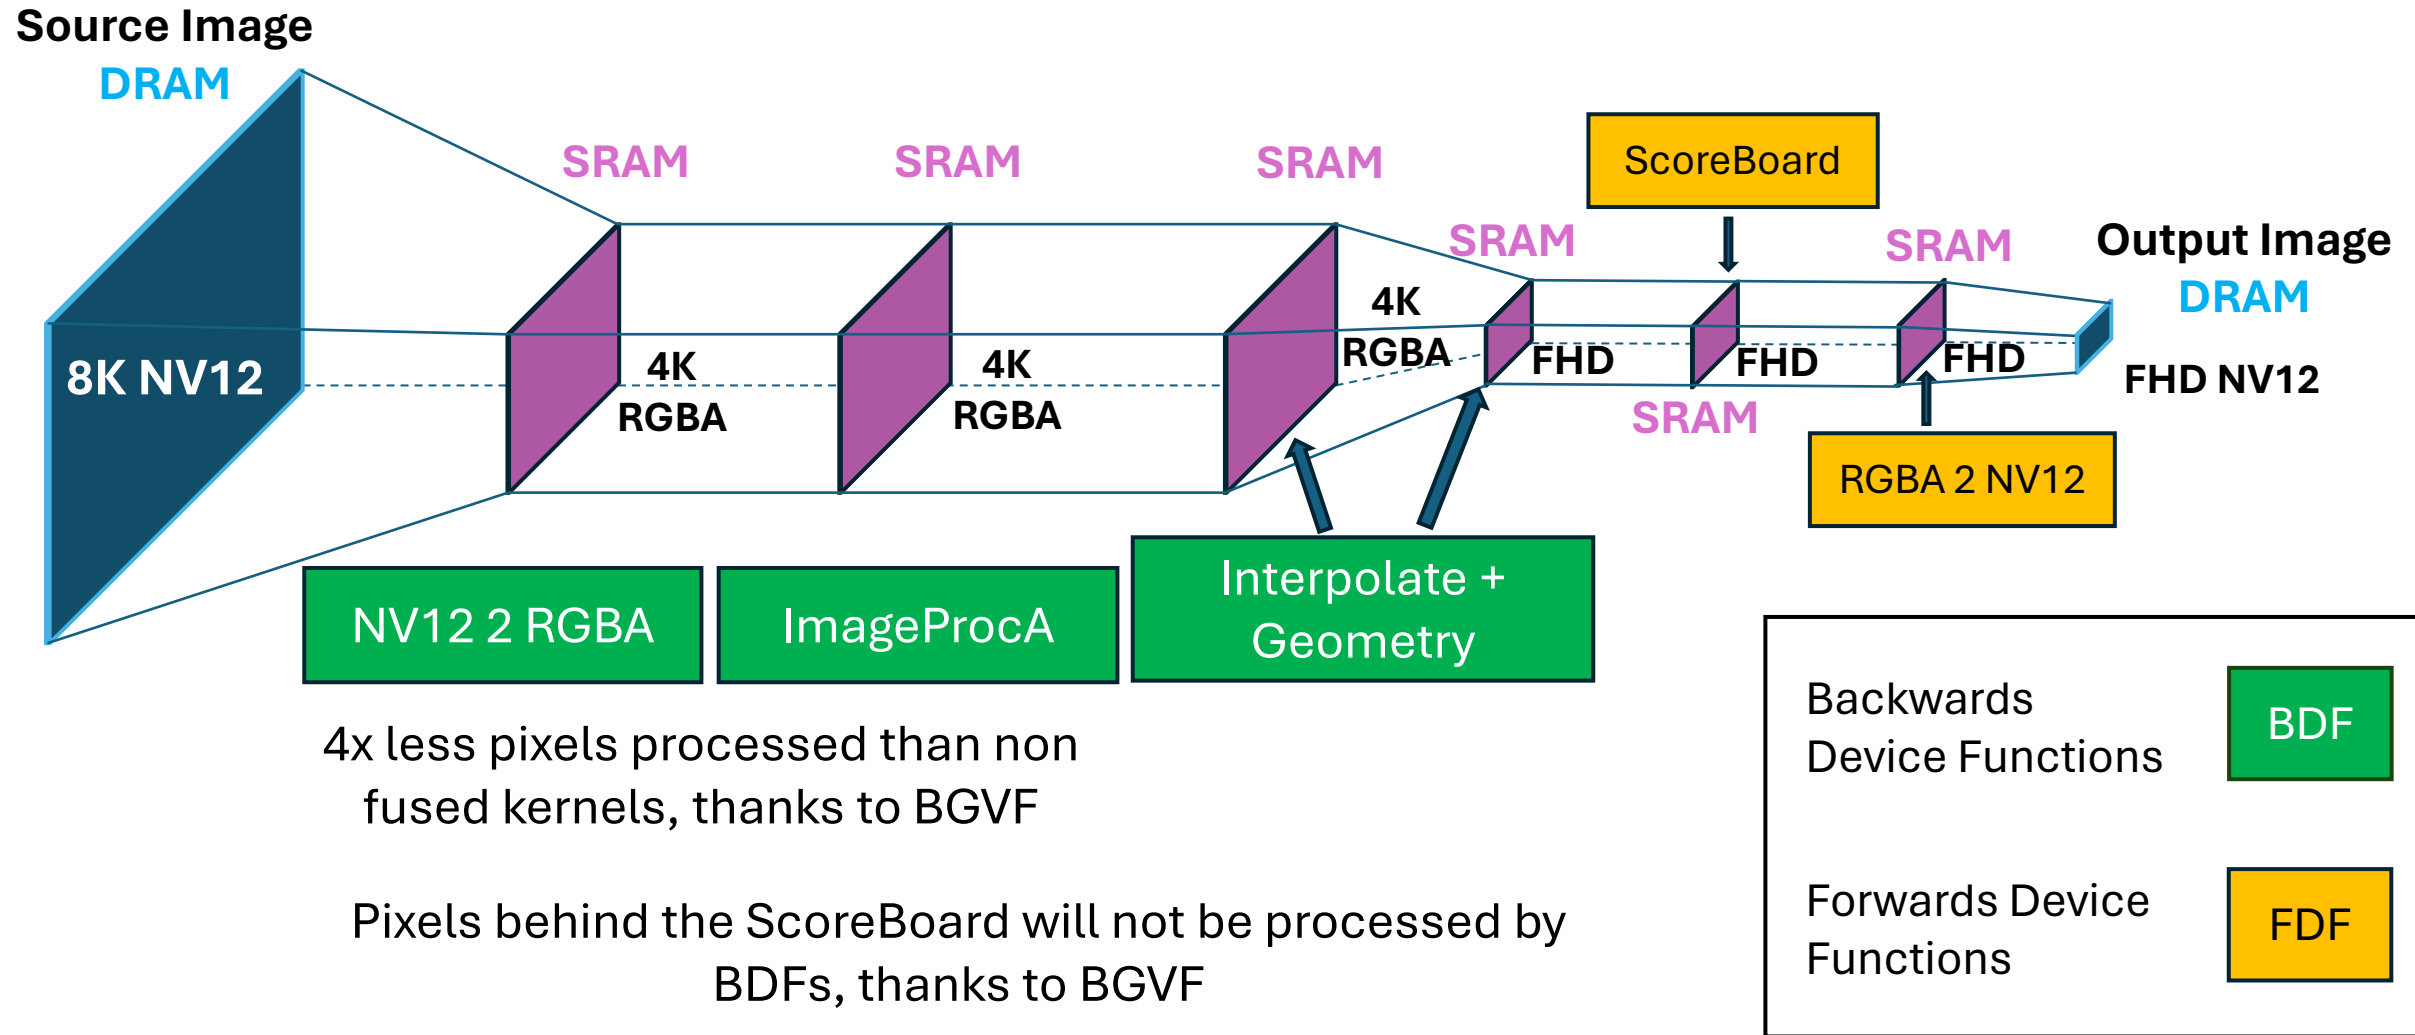
```
using HowToReadAPixel = Binary<ReadYUV<NV12>, ConvertYUVToRGB<NV12, Full, bt709, AddAlpha, float4>, ImageProcA<float4>>;

HowToReadAPixel readDF;
get_params<0>(readDF) = d_nv12Image; // Source 8K image (ReadYUV Device Function)
get_params<2>(readDF) = imgProcAParams; // Parameters required by the ImageProcA Device Function.
// Applying Backwards Generic Vertical Fussion (BGVF)
auto howToInterpolateAPixelDF = resize<HowToReadAPixel, INTER_LINEAR>(readDF.params,
                                                                    Size(d_nv12Image.dims().width, d_nv12Image.dims().height),
                                                                    Size(targetWidth, targeHeight));

// Not present in the OpenSource library, approximated code, more BGVF
auto howToTransformAPixelDF = geometryFuntionBuilder(howToInterpolateAPixelDF, stitchParams);
auto generateOutputImageWithScoreBoardDF = scoreBoardFunctionBuilder(howToTransformAPixelDF, scoreBoardParams);

// Launch a single CUDA kernel, that does everything
executeOperations(stream, generateOutputImageWithScoreBoardDF, // Generic Vertical Fusion after the BGVF
                Binary<ConvertRGBAToYUV<...>>{},
                Write<ChromaSubSampling<NV12>, ...>>{d_outputImage});
```

GPU kernel libraries: graph total fusion



GPU kernel libraries: graph total fusion

Any one imagined a fused neural network with the previous slide?

It will require extra work, but yes, we have many ideas on how to get there, (including reductions, MMA etc...)





GPU kernel libraries

- **Automatic Thread Coarsening:**
 - Compile time detection of the possibility/convenience to apply it or not.
 - Requires zero user intervention (non CUDA programmers)
 - Requires CUDA ninjas to include it or not in their Read and Write Device Functions
 - Around 2x speedups for very memory bound kernels
 - Currently only active for 1 or 2 Byte data types:
 - Bigger types do not seem to give any speedup.
 - Further analysis will look into it.